

A Structure-Aware Hierarchical Graph Neural Network for Multi-Label Scent Profile Prediction in Binary Fragrance Mixtures

Sofia Gabrielle I. Belarmino

Bachelor of Science in Computer Science

Lester Liby V. Ojeñar

Bachelor of Science in Computer Science

Senior Thesis submitted to the faculty of the
Department of Computer Science
College of Computer Studies, Ateneo de Naga University
in partial fulfillment of the requirements for their respective
Bachelor of Science degrees

Project Advisor: Ian Peter L. Lastimosa

Jelly P. Aureus, MS

Jerome G. Flores

Jasper V. Gomez

September 2026

Naga City, Philippines

Keywords: graph neural networks, multi-label classification, set learning

Copyright 2026, Sofia Gabrielle I. Belarmino and Lester Liby V. Ojeñar

The Senior Project entitled

**A Structure-Aware Hierarchical Graph Neural Network for Multi-Label
Scent Profile Prediction in Binary Fragrance Mixtures**

developed by

Sofia Gabrielle I. Belarmino

Bachelor of Science in Computer Science

Lester Liby V. Ojeñar

Bachelor of Science in Computer Science

and submitted in partial fulfillment of the requirements of their respective Bachelor of Science degrees
has been rigorously examined and recommended for approval and acceptance.

Jelly P. Aureus, MS

Panel Member

Date signed: _____

Jerome G. Flores

Panel Member

Date signed: _____

Jasper V. Gomez

Panel Member

Date signed: _____

Ian Peter L. Lastimosa

Project Advisor

Date signed: _____

The Senior Project entitled

**A Structure-Aware Hierarchical Graph Neural Network for Multi-Label
Scent Profile Prediction in Binary Fragrance Mixtures**

developed by

Sofia Gabrielle I. Belarmino

Bachelor of Science in Computer Science

Lester Liby V. Ojeñar

Bachelor of Science in Computer Science

and submitted in partial fulfillment of the requirements of their respective Bachelor of Science degrees is hereby approved and accepted by the Department of Computer Science, College of Computer Studies, Ateneo de Naga University.

Jelly P. Aureus, MS

Chair, Department of Computer Science

Date signed: _____

Marianne A. Tolentino, MS

Dean, College of Computer Studies

Date signed: _____

Declaration of Original Work

We declare that the Senior Project entitled

A Structure-Aware Hierarchical Graph Neural Network for Multi-Label Scent Profile Prediction in Binary Fragrance Mixtures

which we submitted to the faculty of the

Department of Computer Science, Ateneo de Naga University

is our own work. To the best of our knowledge, it does not contain materials published or written by another person, except where due citation and acknowledgement is made in our senior project documentation. The contributions of other people whom we worked with to complete this senior project are explicitly cited and acknowledged in our senior project documentation.

We also declare that the intellectual content of this senior project is the product of our own work. We conceptualized, designed, encoded, and debugged the source code of the core programs in our senior project. The source code of third party APIs and library functions used in our program are explicitly cited and acknowledged in our senior project documentation. Also duly acknowledged are the assistance of others in minor details of editing and reproduction of the documentation.

In our honor, we declare that we did not pass off as our own the work done by another person. We are the only persons who encoded the source code of our software. We understand that we may get a failing mark if the source code of our program is in fact the work of another person.

Sofia Gabrielle I. Belarmino

4 - Bachelor of Science in Computer Science

Lester Liby V. Ojeñar

4 - Bachelor of Science in Computer Science

This declaration is witnessed by:

Ian Peter L. Lastimosa

Project Advisor

A Structure-Aware Hierarchical Graph Neural Network for Multi-Label Scent Profile Prediction in Binary Fragrance Mixtures

by

Sofia Gabrielle I. Belarmino and Lester Liby V. Ojeñar

Project Advisor: Ian Peter L. Lastimosa

Department of Computer Science

ABSTRACT

Predicting the scent of a blend of two aroma chemicals is harder than predicting the scent of one, because a blend does not always smell like the sum of its parts. Graph neural networks (GNNs) now predict the scent of single molecules well, but when extended to blends they usually combine the two molecules by averaging, which assumes the molecules do not affect each other. This study tests whether adding an interaction step improves prediction. Each molecule is encoded by a shared graph attention network, and the two resulting representations are combined in one of three ways: by averaging, by concatenation (a reproduction of a published pairwise model), or by a Set Transformer that adjusts each representation according to its partner before combining them. All three models predict a set of scent descriptors for each blend. The models are trained and tested on a public dataset of two-molecule blends whose labels come from the GoodScents reference catalog. Since these labels are not direct human panel measurements, this study first measures their reliability through agreement between independent catalog sources and label statistics. Each blend descriptor is graded into one of three reliability tiers: Gold, where the blend pairing is independently confirmed by both molecules and the descriptor shows reliable cross-catalog agreement between GoodScents and Leffingwell; Silver, where only one of these two conditions holds; and Bronze, where neither holds. Within a leakage-resistant test split built from UMAP and K-means clustering, the models are compared on the high-confidence (Gold) subset using macro F1, micro F1, PR-AUC, and Hamming loss, with paired bootstrap confidence intervals. Finally, the models’ sensitivity to injected label noise is evaluated, to test whether the interaction-aware model degrades faster or slower than the baselines as label quality declines.

TABLE OF CONTENTS

1	Introduction	1
1.1	Project Context	2
1.2	Purpose and Description	3
1.3	Objectives	5
1.4	Scope and Limitations	6
1.5	Definition of Terms	7
2	Review of Related Literature	11
2.1	Computational Olfaction and Single-Molecule Prediction	11
2.1.1	History of QSOR Modeling	12
2.1.2	The Principal Odor Map	13
2.1.3	Limitations of Single-Molecule Models	14
2.2	Psychophysics of Odor Mixtures	14
2.2.1	Scent Blending Mechanics	14
2.2.2	Elemental vs. Configural Processing	15
2.2.3	Perceptual Linearity vs. Receptor-Level Non-Linearity	16
2.3	Machine Learning on Sets and Chemical Mixtures	18
2.3.1	Permutation Invariance on Sets	18
2.3.2	Precedents in Mixture Modeling	19
2.4	Biological Grounding of Olfactory Interactions	21
2.5	Prior Hierarchical GNNs	22
2.6	Out-of-Distribution Splitting Strategies	23

2.7	Sustainable Fragrance Design and Environmental Considerations	25
2.8	Synthesis and Research Gap	26
2.8.1	Synthesis of Literature	26
2.8.2	Statement of the Research Gap	26
2.8.3	Positioning of the Present Study	27
3	Theoretical and Conceptual Framework	28
3.1	Theoretical Framework	28
3.1.1	Graph-Theoretic Molecular Representation	28
3.1.2	Deep Learning Theory: Graph Attention Networks (GAT)	31
3.1.3	Set Invariance and Set Transformers	32
3.1.4	Readout Expressivity Limits and Symmetry Collapse	33
3.2	Conceptual Framework	34
3.2.1	Input Phase	34
3.2.2	Process Phase 1: Data Preparation and Graph Carving	35
3.2.3	Process Phase 2: Encoding and Interaction	36
3.2.4	Process Phase 3: Classification and Optimization	36
3.2.5	Output Phase	37
3.3	Technical Background	37
3.3.1	Cheminformatics Standard (RDKit): SMILES Parsing and Sanitization . . .	38
3.3.2	Deep Learning Stack: Sparse Graph Convolutions	39
3.3.3	Multi-Label Formulation: Sigmoids and Sparsity Mitigation	40
4	Methodology	41
4.1	Research Design	41
4.2	Data Acquisition and Harmonization Pipeline	42
4.2.1	Dataset Consolidation	42
4.2.2	Vocabulary Standardization	43
4.3	Technical Modeling and Architecture	44
4.3.1	Level 1 GAT Featurization	44
4.3.2	Level 2 Set Transformer Configuration	44
4.3.3	Output Head & Asymmetric Focal Loss:	46

4.4	Leakage-Resistant OOD Splitting and Graph Carving	46
4.4.1	Scaffold Split Decay	46
4.4.2	Morgan-UMAP K-Means Split	47
4.4.3	The Graph-Carving Protocol	48
4.4.4	Survival Vocabulary Filter	48
4.4.5	Interpolative Baseline Sanity Check	49
4.5	Quantitative Metrics and Metric Invariance	49
4.5.1	Multi-Label Metric Invariances	50
4.5.2	Hamming Loss Integration	50
4.6	Statistical Significance Testing	50
4.7	Ablation Study Design	51
4.7.1	Intra-Molecule Encoder (Level 1 GCN Ablation)	52
4.7.2	Inter-Molecule Aggregation (Level 2 Ablation)	52
4.7.3	Attention-Biology Alignment Analysis	53
4.8	Sustainable Substitute-Finder Prototype	54
A	Chemical Feature Vectors	56

LIST OF FIGURES

3.1	Input-Process-Output (IPO) Model.	35
-----	---	----

LIST OF TABLES

A.1	Specification of the 46-Dimensional Atom (Node) Feature Vector (X_i).	57
A.2	Specification of the 10-Dimensional Bond (Edge) Feature Vector (E_{ij}).	57

Chapter 1

Introduction

Perfumers and flavorists often use multiple aroma chemicals and combine them to produce their products [35]. Predicting how combinations of aroma chemical ingredients is difficult, since a blend does not always perform additively, or smell like the sum of its parts [40]. A note in one ingredient may be present but disappear in the blend, and occasionally a note may emerge that neither of ingredients has its own [40, 44]. Existing machine learning models, particularly Graph Neural Networks (GNNs), have performed well on predicting the scent of a single molecule from its structure with accuracy comparable to trained human panelists [25]. However, a limitation of these models is that they were not designed for two molecules at once. When models are extended to blends, the usual approach is to implement naive methods such as mean or sum aggregation, which assumes that the two molecules does not affect each other [44]. Recent human-panel evidence suggests that this assumption often holds, while catalog-labeled blend data suggests it often does not [32]. This study examines whether a model that lets the two molecule summaries interact [52] before combining them predicts the scent labels of two-molecule blends better than one that simply averages them [26].

A further complication is that the blend labels used to train and evaluate such models are not direct measurements of human perception. Public blend datasets, such as the one built by Sisson et al. [40], are compiled from reference-catalog recommendations rather than a trained sensory panel, and their reliability has not been independently measured. This study therefore treats label reliability as a first-order question: before comparing how well an interaction-aware

model predicts blend scent, the reliability of the labels themselves is measured through agreement between independent catalog sources, and the resulting comparisons are reported on a graded, high-confidence subset of the data.

1.1 Project Context

Fragrance communities and scent repositories such as Basenotes and Fragrantica have become a common source of olfactory descriptions. The narratives and scent labels, however, are inconsistent across brands and consumers that create confusion between descriptors that smell alike [6, 49]. Reference catalogs such as The Goodscents Company, and Leffingwell & Associates have been used to train state-of-the-art models [26, 40]. These catalogs are more consistent than community reviews, but their descriptors were not all produced by a controlled sensory panel. This study therefore measures the reliability of the labels it uses before relying on them.

Computational olfaction addresses part of this problem through Geometric Deep Learning by adopting modeling structure-odor relationships (SOR) directly from the nominal chemical structures of chemical compounds rather than relying from written scent descriptions. Graph Neural Networks (GNN) frameworks have become the leading approach of this task, using localized message-passing to learn representations of molecular graphs [31, 40]. A structure to perception study has successfully modeled GNN on expertly curated single-molecule datasets, mapping molecular structure to human perception. The study demonstrated that their model’s prediction of a new molecule is closer to the mean of the 15-member human panel [25]. However, while single-molecule GNNs excel at single-graph classification, single-molecule GNNs are not designed to model a multi-graph set representation. This is the reason why it limits the capability to capture the perceptual and physical dynamics of multi-molecule mixtures, such as masking and emergence [46].

Scents are typically blends with multiple components rather than single compounds. When two aroma chemicals get mixed, the result isn’t just additive; it triggers complex psychophysical and physical interactions. In a study conducted by Sisson et al., they demonstrated that in an unweighted binary mixture, blends exhibited psychophysical distortions, losing an average of 1.2 constituent notes to suppression and gaining 0.2 novel descriptors through configural emergence [40]. Current machine learning models struggle to implement a solution for capturing these biological distortions, largely because they rely on naive linear aggregation, such as mean or sum pooling [54]. Averaging

two molecule embeddings can make structurally different pairs indistinguishable to the model: just as the pairs (2, 8) and (5, 5) share the same average, two chemically distinct blends can average to the same combined representation, discarding exactly the information needed to predict suppression or emergence [44].

1.2 Purpose and Description

This study builds a computational model that predicts blend-level odor descriptors directly from molecular structure, and evaluates that model against a reliability-graded version of its own training data, rather than assuming the underlying labels are trustworthy. This study focuses on predicting the emergent, blend-level odor descriptors of binary (two-molecule) mixtures using hierarchical Graph Neural Networks (GNN) with Set Transformers [26, 40]. Three ways of combining the two molecule representations will be compared: mean pooling, concatenation, and our own model which uses Set Transformer. This study first grades every blend label into a reliability tier and reports how the comparison holds up as label confidence changes.

The level 1 consists of the Intra-Molecular GAT Encoder. A Graph Attention Network (GAT) encodes individual aroma-chemical topologies, producing a molecular representation for each molecule based on its own node attributes and covalent bond features [52]. The encoder uses deep message-passing layers to approximate the mapping from a molecular graph to its isolated, elemental scent descriptors, without yet considering how two molecules will interact.

The level 2 serves as the Inter-Molecular Set Aggregator. This extends the single-molecule GNN architecture into our SA-HGNN (Structure-Aware Hierarchical Graph Neural Network) framework using permutation-invariant mechanisms [26, 54]. These enhancements consist of Self-Attention Blocks (SAB) and Pooling by Multihead Attention (PMA). Within this layer, the PMA module combines the pair of molecule embeddings produced by the encoder into a single, learned blend embedding shaped by interaction between the two molecules, rather than by their individual properties alone. Whether this interactive embedding meaningfully outperforms the averaging and concatenation baselines is evaluated specifically on the high-confidence (Gold-tier) subset of blends, so that any observed advantage cannot be attributed to noise in the underlying labels.

The Multi-Label Classification Head acts as the final output layer. A dense feedforward network maps the combined molecule representation, produced by whichever of the three combination

methods is being evaluated, to independent sigmoid activations over the target scent descriptors. All three configurations are evaluated under the same reliability-tiered protocol, so that differences in reported performance reflect the combination method itself rather than uneven exposure to unreliable labels.

The scent classification model is intended to benefit the following entities:

- **Fragrance Chemists and Perfumers.** A structure-grounded screening tool that models how chemicals interact in a blend can help reduce the cost of trial-and-error cycles during perfume formulation [35].
- **Green Chemical Engineers.** This study contributes a standardized comparison of combination methods for two-molecule mixture prediction, evaluated not only for overall accuracy but for how each method’s performance holds up as label reliability declines. It provides empirical evidence on whether attention-based set decoders overcome the “symmetry collapse” limitation of linear graph readouts [44], and on how set-structured models trained on catalog-derived, noisily-labeled data respond to varying label quality [30].
- **Researchers and Maintainers of Olfactory Reference Catalogs.** The reciprocal-agreement and cross-catalog reliability audit developed in this study is not specific to any one dataset. It offers a reusable procedure for quantifying label trustworthiness in future catalog-derived olfactory datasets, including the very catalogs (GoodScents, Leffingwell) this study draws on, without requiring new human sensory panels to do so.

Research Questions

This study seeks to answer the following questions:

1. How consistent are the two-molecule blend labels, based on (a) cross-catalog reciprocal agreement on each pair and (b) GoodScents-Leffingwell agreement on the parent molecules’ descriptors, and how many blends and descriptors survive when only consistent labels are kept?
2. How does the interaction model compare to the averaging and concatenation baselines on the high-confidence test set, and does this comparison change when training on the full dataset versus high-confidence blends only?

3. How does each model’s performance degrade under increasing label noise, and does the interaction model degrade faster or slower than the baselines?

1.3 Objectives

General Objective

To construct a reliability-graded version of a public two-molecule blend dataset and determine whether a Structure-Aware Hierarchical Graph Neural Network (SA-HGNN), pairing a single-molecule Graph Attention Network (GAT) with a Set Transformer interaction layer, predicts multi-label scent profiles more accurately than averaging and concatenation baselines, and whether that conclusion holds as label reliability varies.

To achieve the general objective, the following specific objectives must be accomplished:

- To preprocess the single-molecule datasets from GoodScents and Leffingwell, and pair them with the binary mixture scent records.
- To measure blend label reliability through reciprocal agreement between catalog records of the same pair and cross-catalog agreement between GoodScents and Leffingwell on the parent molecules’ descriptors, and to assign each blend descriptor a reliability tier (Gold, Silver, or Bronze) on that basis.
- To build the Level 1 Graph Attention Network (GAT) that translates individual chemical structures, specifically their atoms and bonds, into uniform scent vectors, and to pre-train it on individual compounds so the model learns basic structure-to-odor relationships before evaluation on mixtures.
- To implement three ways of combining the two molecule embeddings produced by the Level 1 encoder: averaging, concatenation, and a Level 2 Set Transformer interaction layer using Self-Attention Blocks (SAB) and Pooling by Multihead Attention (PMA), so that the interaction model can be compared directly against the two baselines.
- To implement a Multi-Label Classification Head that maps each combined molecule representation to independent sigmoid activations, predicting the presence of up to 74 catalog-recorded scent notes, with the final count determined by the survival filter after graph carving.

- To partition the dataset into training and test sets using UMAP clustering on molecular fingerprints, so that structurally similar molecules are kept together and no test molecule’s family appears in training.
- To compare the three combination methods on the high-confidence (Gold-tier) test partition, training each on (i) the full dataset and (ii) Gold-tier blends only, using macro F1, micro F1, PR-AUC, and Hamming loss with paired bootstrap confidence intervals.
- To evaluate each model’s sensitivity to label noise by retraining all three combination methods on the Gold-tier training partition with controlled noise injected at 0%, 10%, 20%, and 30%, and comparing the resulting decline in macro F1 across models.

1.4 Scope and Limitations

This study focuses on the design, development, and evaluation of a Structure-Aware Hierarchical Graph Neural Network (SA-HGNN), restricted to pairwise, two-molecule blends. The framework uses a two-level architecture to model blending dynamics directly, rather than relying on simple molecular averaging or naive pooling [40, 44]. The first layer is a Level 1 Intra-Molecular Graph Attention Network (GAT) encoder that processes 2D molecular graphs [52]. The second layer is a Level 2 Inter-Molecular Set Transformer interaction layer, using a Self-Attention Block (SAB) and Pooling by Multihead Attention (PMA) to model non-linear blending dynamics [26, 54]. This interaction layer is evaluated directly against two baselines that combine the same molecule embeddings without an interaction step: mean pooling, and concatenation as reproduced from an existing published pairwise model.

The dataset pipeline merges single-molecule structural and olfactory data from GoodScents and Leffingwell, extracted via the Pyrfume library in Python [6], with a benchmark dataset of catalog-recorded binary chemical blends from Sisson et al.’s study. Model optimization uses a two-phase protocol: pre-training the Level 1 GAT on individual aroma-chemicals to learn structural-olfactory baselines, followed by end-to-end joint fine-tuning on binary mixtures. Optimization uses an Asymmetric Loss function to mitigate the severe label sparsity and class imbalance known to affect multi-label olfactory data [17].

Because the blend labels are compiled from catalog recommendations rather than a controlled sensory panel, their reliability is measured directly rather than assumed. Each blend descriptor is

graded into a reliability tier – Gold, Silver, or Bronze – based on (a) reciprocal agreement between the two catalog records of the same blend pair, and (b) cross-catalog agreement between GoodScents and Leffingwell on the parent molecules’ descriptors, measured using Cohen’s kappa. All three combination methods are evaluated on the resulting high-confidence (Gold-tier) test partition, trained separately on the full dataset and on Gold-tier blends only, to test whether the model comparison holds as label confidence changes.

A leakage-resistant evaluation split is used to assess the model’s out-of-distribution (OOD) generalization [11]. The chemical space is clustered into structurally related families using UMAP dimensionality reduction on 2048-bit Morgan fingerprints (ECFP4), followed by K-Means clustering on the resulting 2D projection [14], so that no test molecule’s structural family appears in training. Performance is compared across the three combination methods using Hamming Loss, Macro and Micro F1-scores, and Area Under the Precision-Recall Curve (PR-AUC), with paired bootstrap confidence intervals. Each combination method’s sensitivity to label noise is additionally evaluated by retraining on the Gold-tier partition with controlled noise injected at 0%, 10%, 20%, and 30%, and comparing the resulting decline in macro F1 across models.

This study is limited to two-molecule blends; concentration or mixing-ratio information is not modeled, and every blend is treated as an unweighted pair. Labels are binary (present or absent) rather than intensity ratings. No new human sensory panel is recruited; label reliability is instead assessed through cross-catalog and reciprocal agreement, as described above. This study does not evaluate whether Level 2 attention weights correspond to documented receptor-level antagonism mechanisms, nor does it develop a standalone ingredient-substitution or screening application; both are noted as potential directions for future work.

1.5 Definition of Terms

Aroma Molecule / Molecular Graph

A volatile organic compound that activates olfactory receptors, represented computationally as an undirected, attributed graph: atoms as nodes and covalent bonds as edges, each described by a short list of chemical properties.

Scent Descriptor (Label)

A word from a fixed vocabulary (e.g., woody, floral, citrus) naming a scent quality. Each molecule

or blend has a present/absent value for every descriptor, drawn from a standardized vocabulary of up to $L = 74$ notes.

Two-Molecule Blend

A combination of two aroma molecules treated as an unordered pair, without a chemical reaction between them: blend $\{A, B\}$ is the same as blend $\{B, A\}$.

Ground Truth (Reference-Catalog Label)

The descriptors recorded for each blend in the dataset used for training and testing, sourced from the GoodScents catalog’s recommended blending partners rather than a controlled sensory panel. This study treats these labels as a proxy for perceived scent and measures their reliability before relying on them.

Reciprocal Agreement

A measure of within-catalog label reliability for a blend pair: whether both molecules independently list each other under the same scent descriptor on their own catalog pages, versus only one molecule making the claim.

Cross-Catalog Agreement (Cohen’s Kappa)

A statistic measuring how consistently two independently maintained catalogs (GoodScents and Leffingwell) assign the same descriptor to the same molecule, correcting for the agreement expected by chance alone.

Reliability Tier (Gold, Silver, Bronze)

A per-descriptor classification of label trustworthiness. A descriptor is Gold if it shows both reciprocal agreement and adequate cross-catalog kappa, Silver if it meets only one condition, and Bronze if it meets neither.

Graph Attention Network (GAT)

A neural network that reads a molecular graph by passing information between connected atoms, learning attention weights that indicate how much each neighboring atom or bond contributes to the molecule’s overall representation.

Molecule Embedding

The fixed-length vector (here, 256-dimensional) that the Level 1 encoder produces to summarize one molecule.

Averaging and Concatenation Baselines

Two linear ways of combining two molecule embeddings into one blend representation: element-wise averaging, or placing them side by side. Because both operations are order-independent by construction, they can only preserve information shared symmetrically between the two molecules, discarding relational information about how the two molecules specifically differ from or relate to one another.

Set Transformer

A network for combining unordered collections of embeddings, built from a Self-Attention Block (SAB), which lets each molecule’s embedding be adjusted according to its partner, and Pooling by Multihead Attention (PMA), which compresses the resulting set into one fixed-length blend embedding. Unlike averaging or concatenation, its output can capture interaction between the two molecules while remaining order-independent.

Leakage-Resistant Evaluation Split

A train/test partitioning protocol in which molecules are first clustered into structurally related families, using UMAP dimensionality reduction on Morgan (ECFP4) fingerprints followed by K -means clustering, so that entire families are held out together. This tests generalization to structurally unfamiliar molecules, rather than allowing near-duplicate molecules to appear on both sides of the split.

Asymmetric Loss (ASL)

A multi-label loss function that down-weights the contribution of easily classified negative labels, preventing the model from being dominated by the many absent descriptors and keeping gradient signal focused on the rare positive ones.

Evaluation Metrics

Macro F1 averages the F1 score across descriptors, giving rare descriptors equal weight; Micro F1 pools predictions globally across all descriptors; PR-AUC measures ranking quality across classification thresholds; Hamming Loss is the fraction of misclassified present/absent decisions across all

descriptors.

Chapter 2

Review of Related Literature

This chapter discusses the relevant studies and systems that support this thesis. First, it tackles seminal research on computational olfaction, followed by the science behind the perception of scents and how different scents behave when mixed. It also examines the application of machine learning to scent prediction when the molecular components in a blend are viewed as a set. Next, the chapter explains the biological architecture beneath these interactions. It reviews hierarchical Graph Neural Networks (GNNs) and their use in the computational modeling of these biological and chemical processes, and presents the literature that justifies the proposed evaluation methodology. Furthermore, it contextualizes the justification for this research, specifically its contribution to sustainable fragrance design, given the ongoing challenges in the domain. This chapter concludes by synthesizing the information from the gathered literature to identify the specific research gap this thesis addresses.

2.1 Computational Olfaction and Single-Molecule Prediction

This first section explores the evolution of computational olfaction from early QSOR modeling with manual feature engineering to more modern graph-based methods. It reviews how these frameworks were used in predicting scent from a single molecule and the limitations of such when applied to the multi-molecule blend prediction task this thesis aims to achieve.

2.1.1 History of QSOR Modeling

Quantitative Structure-Odor Relationship (QSOR) modeling refers to the prediction of the odors and olfactory properties of molecules through the characterization of their structure [13, 50, 55]. It uses the application of chemistry, neuroscience, biology, and psychology to uncover the mechanics behind odor perception, and has long been regarded as one of the key challenges in olfactory research in machine learning [6, 29, 52]. QSOR is especially important in fields such as the fragrance and flavor industries, as it could help streamline the formulation and development of new compounds computationally, as opposed to traditional experimental wet lab methods, which can be both time-consuming and expensive [18, 37].

Early approaches in QSOR involved the manual creation of numerical features that describe and quantify the properties of a molecule [3, 13, 56]. This typically meant either molecular fingerprints, which represent a molecule by encoding the presence or absence of specific substructures into a binary vector [3, 10, 38, 49], or physicochemical descriptors. Examples of fingerprints include Daylight fingerprints, developed by Daylight Chemical Information Systems, which use a path-based featurization approach wherein the connectivity of molecular bonds is analyzed. On the other hand, Morgan or extended-connectivity fingerprints (ECFP) are more modern; instead of tracing linear paths, they take a circular approach, capturing neighboring atoms iteratively within a defined radius [4, 10, 28]. Aside from fingerprints, molecules can also be numerically described using physicochemical descriptors, which characterize molecular structures by their topological, geometric, electronic, and thermodynamic features, computed using software such as Dragon and Mordred [3, 38].

Suh et al., for instance, compared classical molecular descriptors, functional-group fingerprints, and Morgan fingerprints across random forest, XGBoost, and LightGBM on a curated set of 8,681 fragrance compounds, finding that the Morgan-fingerprint models outperformed the descriptor-based ones [43]. Achebouche et al. predicted both the odors of molecules and the olfactory receptors they activate using a set of 5,955 molecules, and found that graph convolutional representations outperformed general molecular descriptors and fingerprints [1]. Saini and Ramanathan demonstrated the use of both Morgan fingerprints and Mordred descriptors, paired with a random forest classifier with multi-label support, to predict odor from molecular structures [37]. Wen et al. also used approximately 5,000 Dragon descriptors with RF, SVM, gradient boosting decision trees, multilayer perceptron, XGBoost, and LightGBM, comparing them against GNN approaches to also predict odor [48].

The common denominator across all of these approaches is that they convert the molecule, specifically its properties, into fixed, handcrafted numerical representations, which are then fed into classical ML models. Recent work such as Suh et al., Achebouché et al., and Wen et al. continues to include descriptor-based classical models, but mainly as baselines to quantify how much fingerprint- and graph-based learning improves upon manual feature engineering [1, 43, 48]. While effective, these methods across both eras require the manual selection of descriptors in advance [4, 25, 56]. This means the model is not actually able to determine on its own which structural features are relevant for the task it is trying to solve. This necessitated the advancement to graph-based deep learning methods, which improve the odor prediction pipeline by teaching models to learn task-relevant features directly from structure [2, 4, 16, 55] which will be discussed in the next section.

2.1.2 The Principal Odor Map

Addressing the aforementioned limitation is Lee et al.’s Principal Odor Map (POM), which maps a molecule’s chemical structure to its smell using Graph Neural Network (GNN) techniques [25]. Rather than evaluating a molecule’s scent using a static representation built from predefined handcrafted descriptors, POM learns directly from the molecule’s structure during training. It does this by treating each molecule as a graph, where nodes and edges represent atoms and chemical bonds, respectively; each node and edge is itself a vector, with nodes storing properties such as atomic number, and edges storing bond properties.

Through message passing, the GNN dynamically assigns weights to each of the molecule’s structural features across multiple passes, producing a numerical representation of the molecule. The final layer predicts the molecule’s specific odor qualities, while the penultimate layer forms the actual Principal Odor Map, which the paper itself describes as a high-dimensional, data-driven representation that is able to preserve human perceptual distances and hierarchies, positioning molecules with similar smells closer together in the generated odor space [25].

The setup was as follows: Lee et al. used approximately 5,000 molecules derived from merging data from the GoodScents and Leffingwell fragrance databases to produce a multi-label output of 138 distinct semantic odor descriptors. It was then validated against a human benchmark, wherein approximately 15 trained panelists rated molecules using a narrowed-down set of 55 odor descriptors. The results show it was able to match the trained panel’s verdicts on every label except for musky scents, and was even able to surpass the median panelist on 30 out of the 55 validation labels (55%)

[25].

This demonstrates the relative effectiveness of GNNs in predicting the scent of a molecule from structure alone. However, despite this significant advancement, these results were derived exclusively from single, unmixed molecules.

2.1.3 Limitations of Single-Molecule Models

Despite GNNs like POM being able to achieve human-like performance in scent prediction, the fact that they only operate on the scale of a single molecule demonstrates a fundamentally limited architecture that fails to reflect the reality that odorants rarely appear in isolation [8, 25, 46, 53]. Real-world scents are rarely composed of single molecules acting on single receptors; rather, they are blends of molecules that can contain individual components numbering up to the hundreds [15, 53]. Consequently, when a method focuses only on a single isolated graph, it is unable to predict new scents that may emerge from the mixture of two or more compounds.

Lee et al. admit this limitation themselves, noting that while POM uses a coordinate system to categorize scents, it alone cannot determine the behavior and outcomes of different molecules when mixed [25]. Moreover, single-molecule coordinate systems do not reflect a biologically plausible architecture, failing to capture the biological complexity of how the mammalian peripheral olfactory system actually operates [7, 19].

2.2 Psychophysics of Odor Mixtures

The unsuitability of single-molecule models for odor prediction in multi-molecular blends signifies the need for a similarly multi-molecular model. This section discusses the groundwork for predicting the behavior of different molecules when mixed, specifically the application of psychophysics to machine learning: how blends deviate from their original descriptors, how the brain processes them, and the concept of linearity in odor perception that will later justify the need for a Set Transformer instead of naive pooling.

2.2.1 Scent Blending Mechanics

A paper released by Sisson et al. introduces a labeled dataset consisting of molecule pairs to measure the divergence of the emerging scent when two molecules are mixed from their original scent labels,

using two metrics: union and intersection. The union score operates under the assumption that the emerging scent will smell like any of the scents either parent molecule has (additive mixing). On the other hand, the intersection score operates under the assumption that the outcome of the two molecules when mixed will only come from the scents the parent molecules share (subtractive mixing) [40].

The Jaccard results show 0.12 for union and 0.24 for intersection [40]. These low scores indicate that neither additive nor subtractive mixing reflects the real behavior exhibited by molecule pairs when mixed, according to the human-reported blend labels. This shows that the naive use of linear math cannot properly replicate the unpredictable behavior that molecule pairs exhibit when mixed.

This non-linearity can be characterized by the phenomena of masking/suppression and emergence, wherein the former describes what happens when a scent present in either parent molecule fails to appear in the blend, while the latter describes what happens when a new scent is perceived that was not present in either parent molecule’s odor descriptors. The results of Sisson et al. show that, on average, 1.2 notes are suppressed and 0.2 notes emerge per blend. Furthermore, an odds-ratio analysis reveals that certain scents, such as "eggy" or "cabbage", were more likely to emerge, while certain scents like "peach" or "fresh" were more likely to be suppressed [40]. This demonstrates that these two phenomena follow a certain learnable pattern, and that the behavior is not merely random, which motivates using a model that can learn these interactions, such as a Set Transformer.

2.2.2 Elemental vs. Configural Processing

Following the examination of the non-linear behavior of the scents of molecule pairs when mixed, this section will provide a sort of basis through the lens of how humans process odors. Studies by Romagny et al. and Coureaud et al. describe two strategies that the human brain uses to process smells, namely configural odour perception and elemental odour perception. With elemental processing, the brain distinguishes the different components that make up a scent. On the other hand, with configural processing, the brain perceives a new scent that is distinct from the original scents that make up its components [9, 36].

To better illustrate their difference, one could imagine two molecules, molecule A and molecule B, which in isolation smell like scent A and scent B respectively. Elemental odour perception would allow a human to smell at least one of, or both, scent A and scent B, while configural odour perception would allow a human to smell a scent that is neither scent A nor scent B.

A real example from Coureaud et al.’s paper describes an experiment wherein molecule A smells like strawberry and molecule B smells like caramel. An A/B ratio of 30:70 was perceived as more pineapple-like, showcasing the brain using configural processing at that ratio; however, at a slight alteration of said ratio, there was a decrease in the perception of pineapple. On the same note, an experiment highlighted in the Romagny et al. paper describes six different molecules that, when mixed together, present a red cordial smell not present in any of its components. However, at a ratio of 16.7% of each component, although the red cordial scent was still present, it was described as less typical of the red cordial scent. The authors hypothesize that an overshadowing effect had occurred within the mixture, whereby one component became perceptually dominant [36]. This effect is consistent with findings where altering component proportions leads to a perceptual shift toward one of the mixture’s individual components, which is more in line with elemental perception. The results of the experiments mentioned in these two papers could imply that the brain activates either configural or elemental processing depending on the concentration, or it could also mean that while the perception mode itself may not change, the typicality of a scent does.

Furthermore, Romagny et al. suggest that a human’s ability to perceive different odors in a single mixture has a ceiling of 3–4, wherein exceeding this might cause the brain to default to configural perception due to the capacity limitations of the brain [36]. Moreover, their paper also finds that different age groups perceive odors differently, with children leaning more toward elemental processing and adults toward configural processing, with teenagers falling in the middle; however, this developmental shift applies only to simple mixtures such as AB.

All of this suggests that a multi-molecule odor prediction model should be able to learn how molecules interact, since emergent scents cannot always be described by the individual components alone. This motivates the Set Transformer used in this study.

2.2.3 Perceptual Linearity vs. Receptor-Level Non-Linearity

While the papers mentioned in sections 2.2.1 and 2.2.2 assert the non-linear behavior of odorants when smelled, this subsection will introduce a paper that could seem to challenge this idea, arguing that the resulting scent from the combination of molecules, despite its perceived “configural” quality at the surface level, is still inherently mathematically predictable and therefore linear.

To prove this, Pellegrino et al. [32] built a quantitative linear model that takes in the unweighted average of an odorant’s descriptors, which are made up of a specialized lexicon of 51 specific but

broad olfactory labels, scored using a rate-all-that-apply system across 432 mixtures and 144 components, as opposed to Sisson’s binary presence/absence labels. Using Pearson correlation on the prediction of overall mixture quality, they were able to yield 0.71, which, when compared against the human benchmark, is not statistically different, proving that a linear model does work for this kind of prediction problem. Furthermore, they also tested famous yet unconventional odorant combinations, such as combining “tomato” and “cucumber” scents together to produce a “cooked fish” smell, yielding a relatively high score of 0.84. This is significant because, even though this kind of relationship would typically be used to describe nonlinearity, their linear model was still able to predict it.

Pellegrino et al.[32] explain this incongruence by claiming that what is causing this perception of non-linearity in odorants is the fact that they are described using simple words only, when in reality, the full olfactory characterization of a molecule holds much more dimensions than that. Taking the previous tomato and cucumber example, taking it at face value makes it seem like there is some nonlinearity when a cooked fish smell seems to develop out of nowhere. In reality, “tomato-smelling” and “cucumber-smelling” are not able to fully capture the complexity of what the scent profile of their respective molecules is really like. Each molecule also contained fish-like elements, but simple words are not enough to describe such occurrences, thus introducing the illusion of nonlinearity. These “hidden” scents exist and have computable values, and since mixing is linear, the hidden fishiness of both molecules was able to combine and thus seem to emerge from nowhere.

However, it is important to note that, despite these discrepancies in how odors are framed, such as whether an odor is configural or elemental, what Pellegrino et al. is trying to argue is that, at the base level, it is all the same when trying to predict the resulting smell, by including in the calculation the “hidden” semantic scent labels a molecule has. A better way to view the different angles that tackle linearity is that both works simply answer the problem differently: Pellegrino et al. uses RATA scores in an unweighted linear average calculation to predict a mixture’s overall quality, while Sisson et al. worked with binary descriptions and used GNNs to produce multi-label descriptors, which is more in line with what this thesis aims to achieve [32, 40]. Pellegrino et al. does acknowledge that, at the receptor level, there are nonlinear interactions that happen, which will be discussed further in a later section.

In any case, in simpler terms, odors could be considered linear when viewed through a continuous scale, while non-linear when viewed through a binary lens. Our thesis specifically takes the angle

of the latter to be able to produce multi-label predictions for odors, similar to Sisson et al.’s [2025] work, therefore justifying our choice of applying a Set Transformer for our prediction task.

2.3 Machine Learning on Sets and Chemical Mixtures

Having established the baseline information on computational olfaction and the behavior and psychophysics of odorant molecules when blended together, this next section discusses the machine learning perspective on the problem, specifically the mathematical foundation underlying the treatment of resulting scent compounds as sets and the relevant prior work on the topic.

2.3.1 Permutation Invariance on Sets

A combination of two or more molecules could be viewed as a mathematical set. An accurate depiction of a fragrance blend of these molecules requires that they be presented in such a way that their positioning in that list is arbitrary, not sequential [40, 46, 55]. For instance, assume the combination of molecule A and molecule B produces scent C. It should be of no relevance whether it is molecule A + B or molecule B + A; it should always produce scent C. For that reason, it is important for any mathematical model used to predict the resultant smell from the combination of different molecules to follow this principle, which is called permutation invariance.

Zaheer et al. [54] established the math of permutation invariance and introduced Deep Sets, a deep learning framework that uses a neural network designed to process datasets in a non-sequential manner for unordered data, which serves a different function from architectures that must process data in order (e.g., molecular compounds vs. protein sequences). This is referenced as the first of two properties they introduced in their study describing permutation invariance, positing that for a function acting upon a set, it must be indifferent to the order in which each element appears, mathematically formulated as $f : 2^X \rightarrow Y$ for any permutation, i.e., for any permutation π :

$$f(\{x_1, \dots, x_M\}) = f(\{x_{\pi(1)}, \dots, x_{\pi(M)}\}).$$

Their second theorem explains how to make the first property possible, stating that the function must be able to individually transform each element ϕ , the results of which are aggregated and then transformed again to ρ , mathematically formulated as decomposable to:

$$\rho \left(\sum_{x \in X} \phi(x) \right),$$

which guarantees the permutation invariance of a set [54].

However, as can be seen in the second theorem, the transformation of each element is done before pooling and is processed in isolation. Therefore, the information about how individual elements may interact with each other is not captured and gets lost, which is pointed out by Lee et al [26]. This is a critical issue in the context of fragrance blends specifically because of the non-linear behavior that odor molecules exhibit. The non-linear behaviors like emergence and masking are not captured if the molecules are treated as having no bearing on each other, which is false.

To bridge this gap, Lee et al. propose the use of a Set Transformer, where instead of processing the elements of a set individually, they apply self-attention mechanisms they call SAB (self-attention block) in the processing of the elements, wherein the different elements are able to learn from each other, such that the extracted features to be used for pooling already contain useful information about the elements' relationship with each other. Moreover, they also developed PMA, or Pooling by Multihead Attention, which replaces the second step in the Deep Sets architecture wherein, instead of fixed sum or mean pooling, PMA also employs self-attention mechanisms in the pooling itself, such that it is able to pick out which of the elements bear more weight for the task at hand. Both of these innovations directly apply to our thesis problem. With the use of SAB, complexities of the interaction between odorant molecules are captured, and made meaningful with the use of PMA.

2.3.2 Precedents in Mixture Modeling

Zaheer et al. [54] provided the groundwork for permutation invariance in sets, while Lee et al. [26] extended it with self-attention, a general set method that later studies applied to molecular mixtures. Plenty of studies have made use of their work in different realms of chemistry.

For example, Qin et al. [33] built the SolvGNN architecture, which aimed to predict the activity coefficient between binary and ternary chemical mixtures by capturing their intermolecular interactions at two levels. First is the atomic-level (local) graph convolutions, which happen in the individual molecules, where features are extracted from each component atom, including the edges that connect them. This structural information is propagated and embedded in preparation for the next level, which is the molecular-level (global) message-passing graph, which operates at a broader scale as the nodes are now the molecules themselves, with the edges representing molecule-to-molecule interactions, such as hydrogen bonding. In this complete graph, message-passing algorithms are employed, using a Gated Recurrent Unit (GRU) to update each node with each pass, to capture

mixture compositions and non-ideal fluid behaviors, with Qin et al. [33] arguing that these dynamic measures are essential to explicitly model the real physical intermolecular forces in mixtures that naive pooling fails to do.

Similarly, Sanchez Medina et al. [38] introduce the Gibbs–Helmholtz Graph Neural Network (GH-GNN), which also predicts activity coefficients of chemical mixtures, but at different temperatures. What differentiates GH-GNN’s architecture from SolvGNN’s is that, instead of having two separate levels for GNN processing, GH-GNN has a single shared GNN, wherein the nodes and edges already contain both local and global vector features in a pair, which get updated over multiple layers of message passing. Afterwards, a global mean pooling operation is done to compress the data from each graph to create a vector embedding, concatenated with the updated global feature to create the final molecular fingerprint, which will then be transferred to another graph called the mixture graph, which maps out the hydrogen-bonding interaction between the solvent and solute molecules. Sanchez Medina et al. [38] defend this design by explicitly building the temperature-dependency relationship into the model’s structure through a Gibbs–Helmholtz-derived thermodynamic expression, rather than relying on a generic GNN to infer it implicitly from data alone. This closely mirrors, and inspired, our thesis’s approach of doing the same shared single-molecule encoder and then a separate mixture-level interaction step, but with odorant pairs instead of solvent and solute.

Lastly, Tom et al. [46] proposed POMMix, a more relevant precedent for our topic, as it is in the field of olfactory perception, aiming to predict the similarity of the scent of two mixtures. Theirs employs a 3-layer architecture, with the first making use of the previously discussed POM, where a GNN processes the atomic bond traits of the molecules along with the scent labels to create a multi-dimensional vector that describes the scent of the molecule. The second layer employs a module they call CHEMIX, which uses a self-attention mechanism across the graph from the first layer, doing set-pooling to produce a permutation-invariant mixture embedding that is able to tell which components contribute to the nonlinear behaviors that odorants exhibit. The last layer of their architecture is the cosine predictive head, which evaluates the similarity of mixtures according to their scent by calculating the scaled cosine distance between their respective embeddings. Their ablation study compared several molecular embedding types including handcrafted RDKit descriptors and pretrained language-model embeddings, combined with both a classical model (XGBoost) and their attention-based CHEMIX module, finding that embeddings derived from POM consistently outperformed the alternatives, reinforcing that structure-aware, attention-based representations capture

mixture behavior more effectively than simpler alternatives.

All three studies tackle three different prediction problems, yet use the same underlying concept of dynamic self-attention mechanisms, as they prove to perform better and better serve the problem of being able to model molecular interaction, as compared to basic aggregation and naive pooling/concatenation, going on to show the importance of using this framework for this paper, which aims to capture intermolecular interactions between odor pairs.

2.4 Biological Grounding of Olfactory Interactions

A scent is smelled when it travels into the air and to our noses, and inside our noses are about 400 different olfactory receptors that get these scents and deliver them to the brain [8, 12, 20, 24]. Building on the computational and machine learning frameworks used to predict chemical mixtures, this section explores the underlying biological mechanisms that govern these olfactory interactions.

In a nutshell, a receptor is like a road, and an odor molecule is like a vehicle. Different vehicles can go into many different specific roads, and many roads can accommodate various specific vehicles as well. When a car goes through the right road, it gets activated [8, 47]. However, there are also times when a vehicle enters a road but does not completely go through it, and then a different vehicle tries to go through the same road, but it cannot pass, because the other vehicle is currently blocking it. The same could happen with odor molecules, and this is called competitive antagonism. In more formal terms, it is when, instead of an activating ligand (agonist) binding on an active receptor or an orthosteric site to produce a biological response, what instead binds is a non-activating ligand (antagonist), which blocks or stops a response from happening [12, 53].

This phenomenon is explored and exemplified by multiple papers. For example, in the paper by Hu et al. [20], the molecule 2,3,5-trimethylpyrazine (TMP), which produces a roasted/nutty note, is supposed to activate receptor OR5K1, but instead is blocked by L-menthol (L-MT), which happens because of analogous binding conformations that prevent co-occupancy of the receptor. A similar thing happens in Fukutani et al. [12], where OR2T-family receptors, which were supposed to be activated by agonist malodors methanethiol (CH_3SH) and hydrogen sulfide (H_2S), instead get blocked by the antagonist β -ionone (a woody/floral ketone). Furthermore, they conducted a nostril experiment wherein: (a) both the antagonist and agonist were inhaled through a single nostril only, versus (b) the antagonist and agonist were inhaled separately through different nostrils. Strikingly,

in test (a), the perceived intensity of the agonist malodor was significantly decreased, compared to test (b), where no suppression occurred, which suggests that suppression in smelling odors happens at the receptor level, as opposed to the brain, and supporting the notion that this is because the two molecules were physically competing for a place at the receptor, rather than signals competing in the brain. This is also supported by another study, wherein the OR51E1 receptor was bound to by the antagonist Florhydral, reducing humans’ perception of the sweaty malodor isovaleric acid (IVA) [21]. This phenomenon is further supported by the Molecule to Olfactory Receptor (M2OR) database, which curates and compiles over 75,000 bioassay experiments and over 51,000 distinct OR-molecule pairs, detailing responsive agonists, non-responsive pairs, and other metadata that aid in the systematic exploration of the combinatorial nature of odorants in olfactory perception [24].

Taken together, these studies show that suppression in odor perception is a physical phenomenon that occurs when molecules compete over the same receptor. What this means for our study is that, just as a receptor’s response to one molecule can be blocked or weakened by the presence of a competing molecule at the same binding site, our Level 2 Set Transformer’s self-attention mechanism allows one molecule’s learned representation to influence how much weight another molecule’s features carry in the final blend embedding. This means that the Set Transformer is not only well-suited for handling molecular sets, but could also be used to model the mechanism of olfactory receptors in the nose, adding a real biological grounding layer to our architecture.

2.5 Prior Hierarchical GNNs

With the biological foundations of this study established, this section will discuss prior hierarchical GNNs in adjacent fields, which mostly operate on single-molecule or atom-to-fragment conditions, and how our proposed hierarchical GNN differs, as it operates between molecules, showing a single-molecule-to-blend hierarchy, which is a distinction that reveals a gap this thesis addresses.

Two formative works in the use of hierarchical graph neural networks for molecular property prediction are Hierarchical Informative Graph Neural Networks (HiGNN) and Hierarchical Molecular Graph Neural Networks (HimGNN) [16, 56]. HiGNN uses an algorithm called BRICS (Breaking of Retrosynthetically Interesting Chemical Substructures), which breaks down molecules into 16 different fragment types to create a graph, which is then combined with its global, non-broken-down counterpart graph to predict the molecule’s property [56]. Similarly, HimGNN makes use

of two graphs, with the first being a global graph as well, composed of nodes as atoms and edges as chemical bonds, and a motif-based graph, where the nodes are motifs, which are the specific molecular substructure patterns that frequently occur, and the edges are the atoms that share them. Both of these hierarchical GNNs are undeniably sophisticated, but they remain single-molecule, atom-to-fragment, meaning their scope is limited to a single molecule. There is other similar work employing GNNs that are non-hierarchical, such as MMGX (Multiple Molecular Graph eXplainable discovery), which employs four different graph topologies for multi-graph representation learning, but is still single-molecule in scope [23]. In olfaction, the previously discussed POMMix remains a rare exception that uses a hierarchical GNN for more than a single molecule; SolvGNN and GH-GNN do so for chemical properties rather than scent. This thus highlights the need for an approach that innovatively uses hierarchical GNNs for processing molecules as compounds.

2.6 Out-of-Distribution Splitting Strategies

This section will be the last technical segment of this chapter, providing a review of the standard data-splitting strategy of scaffold splitting, and an examination of how its application falls short in sparse chemical spaces, which then motivates the need for more rigorous out-of-distribution (OOD) evaluation. This section then grounds and discusses the clustering-based alternative our proposed framework will employ.

A common method to partition data for training, testing, and validation in chemoinformatics and molecular machine learning is scaffold splitting [14, 42]. In this method, molecules are represented by graphs of their structure, and the core framework of these graphs is referred to as their “scaffold” (Bemis–Murcko scaffolds) [51]. Molecules that share the same scaffold are grouped together, and these groups are then distributed among the training, testing, and validation sets to ensure that the test set contains structurally different molecules from the training set, which is a basic form of out-of-distribution (OOD) evaluation, which as opposed to random splitting, simply distributes molecules randomly across training, validation, and test sets [11, 14, 40, 42].

The caveat with this method is that, when the dataset is sparse, it collapses toward random splitting. For instance, similarly structured molecules get fragmented across both the train and test splits, leading to severe data leakage and inflated test scores as scaffolds in standard benchmarks could be populated by only one or two molecules each, undermining the grouping effect the method

depends on [11]. Another problem demonstrated by scaffold splitting is that even molecules that are near-identical (as computed by Tanimoto similarity) [14, 42], even differing by as little as a single atom, can, under this method, end up in opposite splits, as observed by Guo et al. and Steshin et al. [14, 42], who argue that this problem causes models to “memorize” rather than generalize from the data. Because of these problems, scaffold splitting can misleadingly produce scores that are overly optimistic.

To overcome these challenges, Joeres et al. [22] introduced DataSAIL, wherein the data-splitting problem was viewed as a mathematical optimization task; using a heuristic based on clustering and integer linear programming, they aimed to balance minimizing between-group similarity while maximizing within-group similarity, all while maintaining balanced fold sizes and class distributions. On the other hand, Shimakawa et al. [39] propose doing topological pre-clustering by projecting 2048-bit ECFP4 fingerprints onto a 2D map using 2D Uniform Manifold Approximation and Projection (UMAP), then grouping them together using spectral k -means clustering. This was then validated using Leave-One-Cluster-Out Cross-Validation (LOCOCV), where clusters are left out in rotation for training and testing, such that what the model was trained on does not show up in what it is tested on, meaning the model actually learns, rather than simply memorizing the data it was trained on, which is important for OOD evaluation.

However, while these methods may work well for single-molecule data splitting, they would be insufficient for multi-molecular fragrance mixtures, because an individual molecule can be a component of varying other molecules in different blends, meaning a molecule seen during training could reappear as a component in a different, unseen test blend, leaking information about that molecule across the split, causing the performance metrics to inflate. Therefore, Sisson et al. [40] propose “graph carving”, using a meta-graph to describe the pairwise/multi-molecule blend dataset, where each node represents an individual molecular graph and each edge represents a tested blend combination. Edges connecting a molecule from one partition to a molecule in another are discarded to guarantee no component overlap between splits. In their original implementation, this partitioning was done by randomly assigning molecules to each side and repeating until every label had at least one training and one test example, since a true scaffold-based split was not compatible with their sparse label distribution.

Building on this, our proposed framework instead applies the graph-carving protocol on top of the UMAP + k -means clustering discussed above, rather than random or scaffold-based partition-

ing. By applying this combined methodology to our merged GoodScents–Leffingwell–Sisson dataset, clustering molecules via UMAP and k -means, then carving the resulting blend meta-graph along those cluster boundaries, our proposed framework is evaluated on a strictly out-of-distribution test set.

2.7 Sustainable Fragrance Design and Environmental Considerations

As the architectural and methodological considerations have been discussed, this section positions this study as motivated by real-world problems which include the growing regulatory and environmental pressures affecting the fragrance industry, and the role computational tools can play in addressing them.

The development of fragrances is an intricate, costly, and time-consuming process, relying heavily on expensive, empirical trial-and-error methods that can take up to three years and cost up to \$50,000 per kilogram for perfumes specifically [35]. Beyond these hurdles, safety and environmental regulations facing the fragrance and flavor industry demand sustainable ingredient substitution [18]. Stricter global safety mandates, such as the European Union’s REACH regulations and evolving International Fragrance Association (IFRA) Standards, have caused legacy aroma-chemicals to become obsolete due to issues in material biodegradability, skin sensitization, and environmental toxicity [30, 46]. In response, initiatives such as the Safe and Sustainable by Design (SSbD) framework under the EU Chemicals Strategy for Sustainability were implemented to promote chemical safety and circularity across the product lifecycle.

To address these challenges, researchers have turned to Computer-Aided Molecular Design (CAMD) and Quantitative Structure-Property Relationship (QSPR) frameworks. However, as it stands, the selection guides used in traditional green substitution frameworks are static and contain fewer than 200 solvents, as expanding these guides would require collecting more than 30 distinct physical properties per molecule—a process that demands laborious laboratory testing [10]. Through the use of GNNs, which can dynamically generate descriptors for odorant molecules without the need for a priori engineering, in-silico pre-screening for candidate compounds can expand the searchable chemical landscape and significantly reduce laboratory trial-and-error processes [4, 10].

However, a candidate substitute would only be qualified for commercial fragrance formulation

if it meets two essential conditions: physical compatibility and smell potency [10, 27]. Physical compatibility refers to how smoothly the substitute dissolves in the base it is mixed into, and is evaluated using Hansen Solubility Parameters, where the Relative Energy Difference (RED) must be less than 1 [10, 35]. Potency matching, on the other hand, refers to predicting Odor Detection Thresholds (ODT) and Odor Activity Values (OAV), such that the target scent strength is achieved while balancing the safety and cost-effectiveness of the substitute’s concentration [27, 35]. These two requirements then serve as the basis for our Sustainable Substitute-Finder, which uses nearest-neighbor retrieval in the Level 1 GAT embedding space to find non-toxic, eco-friendly alternatives that preserve the original scent profile.

2.8 Synthesis and Research Gap

2.8.1 Synthesis of Literature

Current models in computational olfaction can successfully map single-molecule structures to their appropriate sensory descriptors [25]. However, because these compounds are evaluated in isolation, nonlinear psychophysical interactions such as masking and emergence—which determine what the final resulting scent would smell like when multiple odorant molecules are mixed—are not taken into account [20, 36, 40].

From an architectural standpoint, permutation-invariant set processing must be used to predict the behavior of multi-molecule mixtures [54, 26]. Standard linear pooling (mean/sum) methods are incapable of handling non-additive interactions due to symmetry collapse [44]. Current hierarchical GNNs, on the other hand, only consider intra-molecular fragments and fail to capture intermolecular blend interactions [16, 56].

Furthermore, standard scaffold splits suffer from severe data leakage at the component level in mixture spaces [11]. As a result, continuous manifold clustering must be used in conjunction with meta-graph carving to guarantee leak-free evaluation [39, 40].

2.8.2 Statement of the Research Gap

There is a lack of research evaluating an explicit, permutation-invariant attention mechanism (Set Transformer) against naive linear vector pooling on a standardized, leakage-resistant, out-of-distribution

(OOD) benchmark for multi-label fragrance blend prediction.

2.8.3 Positioning of the Present Study

To close this gap, this study presents a Structure-Aware Hierarchical Graph Neural Network (SA-HGNN) that applies a weight-shared Level 1 Graph Attention Network (GAT) to encode information on a per-molecule basis, and makes use of an explicit Level 2 Set Transformer (SAB + PMA) to handle the interactions between molecules, thereby allowing the model to computationally capture intermolecular interactions.

This framework is evaluated on a leakage-resistant Morgan-UMAP k -means graph-carved benchmark to quantify the value of explicit set attention over naive pooling, while establishing a green ingredient substitution engine under UN SDG 12 (Responsible Consumption and Production).

Chapter 3

Theoretical and Conceptual Framework

3.1 Theoretical Framework

This section presents the theory behind the SA-HGNN architecture and the label reliability analysis. It first describes how each aroma molecule is represented as a graph of atoms and bonds. It then explains Graph Attention Networks (GAT), which encode each molecule, followed by set invariance and Set Transformers, which combine the two molecules of a blend. The fourth part explains why simple averaging can lose information about how the two molecules relate to each other, which motivates testing an interaction layer against averaging and concatenation baselines. Because this argument describes a worst case and not an observed result, whether the interaction layer improves prediction is treated as an empirical question. The final part covers the measures used to assess label reliability: reciprocal agreement, Cohen’s kappa, and learning from noisy labels.

3.1.1 Graph-Theoretic Molecular Representation

Each aroma molecules is modeled as an undirected, attributed molecular graph in which atoms are nodes and bonds are edges, and create a mathematical representation where atomic and bond features are updated via message passing between the connected atoms [52]. This would allow the Graph Attention Network (GAT) to extract structural scent profiles regardless of atom ordering.

In cheminformatics, chemical compounds are often written as one-dimensional character strings, such as text-based Simplified Molecular Input Line Entry System (SMILES) strings. However, these text-based strings are syntax-sensitive. A single molecule can be written as many as single valid string depending on the starting atoms, and small edits can produce a completely different molecule [31, 49]. Representing a molecule as a graph captures its bond connectivity without depending on the order in which atoms are listed, which matches how the Graph Attention Network passes information between bonded atoms.

Mathematical Formalism

Each aroma molecule will be defined as an undirected, attributed graph $G = (V, E)$, where:

- V represents the finite set of $N = |V|$ nodes, corresponding to the non-hydrogen atoms in the molecule. Hydrogen will be treated as implicit properties embedded directly within the feature vectors of their parent heavy atoms, rather than independent nodes, which will reduce the graph complexity and computational overhead [4].
- E represents the set of $M = |E|$ undirected edges. An edge $(v_i, v_j) \in E$ represents a covalent bond between heavy atoms v_i and v_j . Because a bond is a mutual sharing of electrons between two atoms, it has no direction: $(v_i, v_j) \in E$ if and only if $(v_j, v_i) \in E$ [4].

The connectivity of the molecule is described by a symmetric binary adjacency matrix $A \in \{0, 1\}^{N \times N}$, where $A_{ij} = 1$ if a bond exists between atoms v_i and v_j and $A_{ij} = 0$ otherwise [56]. In implementation, the graph is stored in a sparse format instead of a dense matrix (Section 3.3.2). Each node and edge also carries a feature vector:

- Node vector $\mathbf{x}_i \in \mathbb{R}^{46}$, which is the initial state $\mathbf{h}_i^{(0)}$ used by the GAT. It records the atom’s element, formal charge, number of bonds, valence, hybridization, aromaticity, attached hydrogens, ring membership and ring size, and chirality, in 9, 1, 6, 7, 6, 1, 5, 6+1, and 4 dimensions respectively, for a total of 46 (Table A.1) [41, 56].
- Edge vector $\mathbf{e}_{ij} \in \mathbb{R}^{10}$, which records the bond’s type, conjugation, ring membership, and stereochemical configuration (Table A.2) [56].

Olfactory Domain Mapping

Some structural details are known to affect scent, including chirality (the two enantiomers of carvone, for example, smell of spearmint and caraway), ring structure, and aromaticity. The node and edge features encode these details so that the model has access to them. Features are computed with the RDKit cheminformatics library [6, 41, 25], using feature types common in molecular and odor-prediction models [25, 41]. The features describe individual atoms and bonds. Larger structures such as functional groups are not encoded directly; the GAT can learn to recognize them by combining information across neighboring atoms. The exact composition of the 46-dimensional node feature vector and the 10-dimensional edge feature vector is given in Appendix A.

Analytical and Methodological Justification

Aroma molecules modeled as graph structures instead of physical descriptors or using the text-based SMILES strings provide key advantages:

- **Order Independence:** The encoder sees the molecule as one connected structure as a sequential list but as a whole, so the output does not depend on the sequence in which atoms are listed. This concerns atom order within one molecule and is separate from the order independence of the two-molecule blend discussed in Section 3.1.3. This ensures that it predicts scents no matter how the atoms are arranged [31].
- **Localized Feature Preservation:** Specific clusters of atoms, called “osmophores,” are thought to be responsible for particular odor qualities [41]. Since each atom carries its own features and the GAT aggregates information from bonded neighbors, the encoder can learn patterns in local chemical environments before combining them into a molecule-level embedding [52, 25]. The model is not given receptor data, so it does not identify receptor-binding sites.
- **Isomers:** Molecules with similar atoms but different bonding—*isomers*—have different graphs and are therefore distinguished. Molecules that have the exact same atoms and bondings called *stereoisomers* are distinguished if explicit attributes—chirality and double-bond configuration—are added into its features, and 2D graph does not capture 3D shape or conformation [4, 49, 55].

3.1.2 Deep Learning Theory: Graph Attention Networks (GAT)

The model uses a Graph Attention Network (GAT) to map each molecular graph to a numerical vector. Atoms are the nodes and the bonds connecting them are the edges. Instead of combining a neighboring atom’s information with fixed weights set by the graph structure, a GAT learns a weight for each neighbor, so the encoder can give more influence to some neighboring atoms than others when building its vector representation [25, 52]. This is loosely analogous to olfactory receptors responding more strongly to some structural features than to others, but the learned weights are model parameters and are not evidence of receptor behavior or chemical importance.

Let $\mathbf{h}_i^{(0)} = \mathbf{x}_i$ denote the initial 46-dimensional feature vector of atom v_i . A learnable weight matrix $\mathbf{W}$ projects the features of each atom into a hidden space of size $F' = 256$. To score how much atom v_i should attend to a neighbor v_j , a single-layer feedforward network with parameters $\mathbf{a}$ evaluates the concatenated projected features of the two atoms and the projected 10-dimensional bond feature vector $\mathbf{e}_{ij}$ between them [52, 56]:

$$\tilde{e}_{ij} = \text{LeakyReLU}\left(\mathbf{a}^T[\mathbf{W}\mathbf{h}_i \parallel \mathbf{W}\mathbf{h}_j \parallel \mathbf{W}_e\mathbf{e}_{ij}]\right), \quad (3.1)$$

where $\parallel$ denotes concatenation. Including $\mathbf{e}_{ij}$ lets the score depend on the bond type, conjugation, ring membership, and stereochemical configuration. The LeakyReLU, with a negative slope of 0.2, keeps a small gradient for negative scores so that they do not become stuck at zero [52].

Atoms have different numbers of neighbors, so the raw scores are normalized across the neighborhood $\mathcal{N}(i)$ of atom v_i with a softmax, so that the weights for each atom sum to one [52, 4]:

$$\alpha_{ij} = \text{Softmax}_j(\tilde{e}_{ij}) = \frac{\exp(\tilde{e}_{ij})}{\sum_{k \in \mathcal{N}(i)} \exp(\tilde{e}_{ik})}. \quad (3.2)$$

The normalized weights α_{ij} scale the projected neighbor features, and the Exponential Linear Unit (ELU) activation is applied to give the updated representation of atom v_i [52]:

$$\mathbf{h}_i^{(l+1)} = \text{ELU}\left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij} \mathbf{W}\mathbf{h}_j^{(l)}\right). \quad (3.3)$$

To stabilize learning, the encoder runs $K = 4$ independent attention heads in parallel. In the hidden layers their outputs are concatenated, and in the final layer they are averaged [52]. Global mean pooling over the atoms of the molecule then gives the final 256-dimensional molecular embedding.

3.1.3 Set Invariance and Set Transformers

Mixing two fragrance ingredients regardless of their order, the resulting scent will smell exactly the same. Mathematically, mixing Molecule A and Molecule B will have the exact same scent as mixing Molecule B and Molecule A. The blend follows the commutative property, which means that it behaves as an unordered set where the order of input does not change the output. To ensure that the blend prediction framework of this study follows the biological and chemical realism, it must be permutation-invariant [54, 40].

The Deep Sets Formulation

Deep Sets was established to solve this order-independence problem. It was mathematically proved that this framework guarantees permutation invariance by processing each set element through an identical neural network transformation before aggregating their representations using a symmetric function, such as summation and averaging [54]. Formally, a function $f(X)$ operating on a set of molecular embeddings X is defined as:

$$f(X) = \rho \left(\sum_{x \in X} \phi(x) \right). \quad (3.4)$$

The function ϕ represents an encoder network that projects each individual molecular embedding x to an intermediate latent space, while the sum operator aggregates these individual vectors [54]. This function is the level 1 Graph Attention Network (GAT) encoder in this study, turning each separate molecule into a 256-dimensional vector. This formulation guarantees order-invariance and is commutative because it treats each molecule of the same dimension. However, the translator ϕ evaluates Molecules A and B in isolation. Deep Sets assumes that these molecules do not affect each other. Real-world olfactory mixtures are complex and exhibit psychophysical phenomena such as scent masking and sensory emergence. The model is unable to capture the non-linear chemical interactions [40, 44].

Set Transformers and Attention-Based Pooling

The complex interaction between the molecules happens at the receptor level, which GAT alone cannot capture. To model these blending effects while maintaining permutation invariance, the model adapts the Set Transformer, which will allow molecular representations to interact dynamically

[26]. The Set Transformer employs two core blocks. The Set Attention Block (SAB) first allows molecule embeddings in a set to cross-attend to one another, enabling each molecule’s representation to be conditioned on its blend partner before pooling. This cross-molecule interaction captures psychophysical phenomena such as masking and emergence that a single encoder cannot model in isolation.

The second block is Pooling by Multihead Attention (PMA):

$$\text{PMA}_k(X) = \text{MAB}(S, \text{rFF}(X)). \quad (3.5)$$

This process begins by taking the molecular embeddings $\mathbf{X}$ and passing them through a row-wise feedforward network that prepares each molecule’s data for the final pooling calculations. The PMA block introduces a set of k learnable seed vectors, denoted as (S) . These probes are refined during training to actively extract the dominant features of the scent blend. The Multihead Attention Block (MAB) lets these seed vectors (S) attend to the polished molecule embeddings $(\text{rFF}(X))$, compressing the variable-size set into a fixed-size blend representation. This attention-based pooling captures the receptor-level interactions, providing a biologically and chemically realistic foundation for scent blend prediction [40, 26].

3.1.4 Readout Expressivity Limits and Symmetry Collapse

The chemical embeddings must be merged into a single representation to capture the exact blend of the molecules. Standard Graph Neural Network (GNN) designs are usually simple when merging vectors. Standard GNNs typically use linear permutation-invariant pooling, such as global sum or mean pooling [54]. Developments in representation theory prove that these impose a severe, encoder-independent expressivity barrier known as symmetry collapse [44]. The mathematical bottleneck limits the ability of standard GNN models from detecting psychophysical phenomena and non-linear chemical relationships, such as masking and emergence [40]. This argument concerns the mixture-level readout, where two molecule embeddings are combined; the within-molecule mean-pooling readout in Level 1 follows common practice in molecular GNNs and is not the focus of this study.

The Mathematics of Symmetry Collapse

Consider a molecular graph G whose set of node embeddings M is in a d -dimensional representation space $\mathcal{U}$. The structural symmetries of the graph are expressed by its automorphism group $\mathcal{H}$. It

consists of a list of all valid permutation node indices, whether flipped, rotated, or swapped; the molecules don’t change the way they are. When we apply a linear, $\mathcal{H}$ -invariant readout function f to compile these embeddings, the function is mathematically forced to factor through the group-averaging Reynolds operator (p_{avg}), defined as:

$$p_{\text{avg}} = \frac{1}{|\mathcal{H}|} \sum_{h \in \mathcal{H}} \rho(h). \quad (3.6)$$

The Reynolds operator projects the entire representation space $\mathcal{U}$ onto its fixed, trivial subspace $\mathcal{U}_1$ [44]. Consequently, any linear $\mathcal{H}$ -invariant map can be decomposed into the composition of a downstream linear map and the group-averaging projection:

$$f = f' \circ p_{\text{avg}}. \quad (3.7)$$

This factorization proves that a linear pooling readouts can only see what survives in $\mathcal{U}_1$. It completely projects away and discards the non-trivial, symmetry-aware components ($\mathcal{U}_2 \oplus \dots \oplus \mathcal{U}_m$) that capture node-to-node variation [44]. Even when the GAT encoder produces node embeddings that are genuinely sensitive to structural differences between molecules, none of that survives a linear readout at the mixture level. Molecular combinations with the same sum or the same average will smell and look alike after the structural information of the molecules is crushed, making it harder for standard GNNs to distinguish molecular combinations [44, 40].

3.2 Conceptual Framework

The conceptual framework of this study operationalizes the theoretical formulations into a systematic computational pipeline structured around the Input-Process-Output (IPO) model. Figure 3.1 outlines the structural interaction between input data representations, the two-level hierarchical deep learning architecture, and the downstream analytical and practical deliverables.

3.2.1 Input Phase

The computational data pipeline starts with the acquisition, curation, and gathering of data from olfactory databases. The raw SMILES strings are extracted from the GoodScents and Leffingwell repositories accessed via the Pyfrume Library in Python [6]. These raw chemical blueprints are then paired with expertly annotated odor tags to establish baseline inputs required from the pre-training




Figure 3.1: Input-Process-Output (IPO) Model.

phases. The binary mixture profiles are retrieved from the Sisson et al. (2025) benchmark dataset [40]. The retrieved benchmark dataset provides standardized chemical pair identifiers (CIDs) and multi-label scent descriptors that will help the network to learn and predict the combined scent of two scent blends.

3.2.2 Process Phase 1: Data Preparation and Graph Carving

To raw chemical input will first pass through a standardization pipeline. Salts and other non-fragrance components are removed using RDKit’s SaltRemover module, leaving the main molecule. Compounds containing elements other than C, H, O, N, S, F, Cl, Br, and I are dropped, since these fall outside the range typical of aroma chemicals [31]. The molecular weight should be between 20 and 600 g/mol. Below this range is considered irrelevant, while above the range is considered too heavy to be perceived as a scent at all [25]. The surviving SMILES strings are canonicalized using RDKit. The scent descriptors text from the source databases is standardized by merging redundant terms and eliminating descriptors that appear rarely to learn from [8]. The blend labels are then audited for reliability before any model is trained. Reciprocal agreement is measured for each blend pair, and cross-catalog agreement between GoodScents and Leffingwell is measured on the

parent molecules using Cohen’s kappa. Each blend descriptor is assigned a Gold, Silver, or Bronze tier based on these two measurements, with the tier rules fixed in advance. To test how well the network handles unfamiliar compounds, molecules are grouped into structural families by UMAP and K-Means on 2048-bit Morgan fingerprints [14]. The Graph-Carving algorithm then divides the mixture meta-graph along family clusters. This guarantees that validation and testing pairs do not contain any molecule encountered during training [11].

3.2.3 Process Phase 2: Encoding and Interaction

The second phase coconverts the flat, two-dimensional 2D molecular graphs into continuous high-dimensional vector spaces and models non-linear interactions between paired compounds. In the encoder, individual atoms and bonds are processed as graph nodes and edges [52]. Since the Level 1 GAT shares its weights, it encodes both molecules under the same chemical rules, which then guarantees symmetry [38]. After encoding and updating atomic neighborhoods, global mean pooling condenses these vector spaces into a single 256-dimensional embedding for every compound [38, 52].

The two embeddings are then combined into a single representation using one of three methods, which are compared against each other mean pooling; concatenation, reproduced from an existing published pairwise model [40]; or the Level 2 Set Transformer, in which each embedding is adjusted according to its molecular before the two are pooled into a single 256-dimensional vector [26]. The first two combine the embeddings through a fixed operation, while the Set Transformer lets the two molecules interact. Whether that interaction improves prediction is the central question of the study.

3.2.4 Process Phase 3: Classification and Optimization

Once the Level 2 yields the 256-dimensional blend vector spaces, it will then be translated into multi-label scent predictions. This Multi-Layer Perceptron (MLP) head contracts sequentially from 256 to 128 and ends with L individual outputs to match the target scent classes [40]. Because scent profiles are inherently multi-label, the classification layer utilizes independent sigmoid activation functions [37]. These activations compute individual probability scores for the joint presence or absence of each unique target scent descriptor. The system is trained end-to-end for network optimization. This means feedback from prediction errors flows backward through all layers simultaneously, fine-tuning both the Level 1 molecular encoders and the Level 2 mixture blocks guided with Asymmetric Focal

Loss [17]. Each combination method is trained under two data conditions, on all reliability tiers and on Gold-tier blends only, and, for the noise sensitivity test, on Gold-tier blends with a controlled fraction of labels corrupted. Olfactory datasets are highly sparse, meaning a single scent profile contains only a few active descriptors alongside dozens of inactive ones (zeros). Under standard training rules, the model would quickly learn that it can achieve over 97% accuracy simply by predicting “no” for every label. The asymmetric loss solves this class imbalance by ignoring these obvious negatives and forcing the network’s mathematical focus onto the rare, active scent notes that are difficult to predict [17].

3.2.5 Output Phase

The final phase outputs a multi-label binary vector of predictions across all L target scent descriptors which is based on the probability given to each descriptor. A validation module evaluates the model’s performance on these predictions under specified conditions. It follows and calculates metrics such as macro- and micro-averaged F1-scores, Hamming loss, and the Precision-Recall Area Under the Curve (PR-AUC) [3, 45, 48]. These metrics measure the model’s performance, ensuring that it generalizes well to new chemical structures. The three combination methods are compared on the high-confidence (Gold-tier) test set, and also on the full test set for transparency. Each method is trained on all reliability tiers and on Gold-tier blends only, and differences between methods are assessed using paired bootstrap confidence intervals. To test how each method responds to unreliable labels, the models are also retrained on Gold-tier blends with a controlled fraction of labels corrupted, and the decline in macro F1 is compared across methods.

The study additionally reports the results of the label reliability audit: reciprocal agreement rates, cross-catalog agreement by descriptor, and the number of blends and descriptors in each reliability tier.

3.3 Technical Background

Mapping molecular structure onto multi-label olfactory predictions requires stitching together a few different layers of infrastructure: cheminformatics standards to make sure the molecules are valid, a graph learning framework to actually process them, and a loss function suited to how sparse the labels are [31]. This section walks through each piece—the parsing standards, the sparse graph

representation, and the loss objective—and explains why each choice was necessary rather than incidental.

3.3.1 Cheminformatics Standard (RDKit): SMILES Parsing and Sanitization

This study utilizes the text-based Simplified Molecular Input Line Entry System (SMILES) strings used in molecular structures. SMILES strings are canonicalized and parsed into undirected, attributed molecular graphs. It uses the open-source RDKit library in Python [31]. Using these strings is convenient because they are compact and suitable for storage. However, this compactness comes at a cost: they are structurally sensitive, prone to syntax errors, and lack explicit physical geometry [49]. Before anything reaches the model, every string goes through RDKit’s sanitization pipeline, which runs four checks in sequence:

- Kekulization checks that the alternating single and double bonds within conjugated or aromatic rings are represented in a chemically consistent manner, resolving ambiguous resonance structures.
- Aromaticity Perception evaluates the coordinate and topological structure of each atom to identify aromatic rings, dynamically reassigning aromatic bond attributes to relevant edges.
- Valence Checking checks every atom, specifically heavy atoms—carbon, nitrogen, oxygen, sulfur, and halogen to ensure it does not exceed its physical valency limits. This rejects syntactically corrupted SMILES strings or physically impossible hypervalent structures.
- Ring Perception, which identifies and indexes ring systems, information the later convolutions depend on.

By enforcing these rules, it ensures that what the network receives is something that is already known to be chemically valid, rather than forcing the neural network to infer the laws of organic valency implicitly from training data [4, 56]. The sanitization pipeline acts as an explicit structural filter for the model. This pipeline lets the Level 1 Graph Attention Network (GAT) operate exclusively on sparse, topologically accurate representations that map directly to real-world volatile compounds [25, 52].

3.3.2 Deep Learning Stack (PyTorch / PyTorch Geometric): Sparse Graph Convolutions

After sanitization, each molecule is formatted as an undirected graph $G = (V, E)$, where atoms represent nodes (V) and covalent bonds represent edges (E). A dense $N \times N$ adjacency matrix could represent this connectivity, but for molecules, where most atoms bond to only a few neighbors, that approach wastes memory, scaling at $\mathcal{O}(N^2)$ for structures that are inherently sparse. The architecture instead runs on PyTorch Geometric, which handles molecular graphs as sparse attributed objects rather than dense arrays [31, 49].

Topology is stored in Coordinate (COO) format: a sparse two-row index matrix C of size $2 \times 2E$, where the first row lists source node indices and the second lists receiver node indices [4]. Since chemical bonds are physical and undirected, every source node is also a receiver, which is why $2E$ entries are needed rather than E — each bond has to be represented in both directions for message-passing to work [4]. Node states are then updated by running sparse convolutions over this index matrix, a process that breaks down into three phases at layer l .

- **The Message Phase:** For each neighbor v_j connected to a central node v_i , the network builds a directed message from their hidden states plus the bond features e_{ij} connecting them [56]:

$$\mathbf{m}_{ij}^{(l)} = \text{Message}^{(l)}(\mathbf{h}_i^{(l)}, \mathbf{h}_j^{(l)}, \mathbf{e}_{ij}). \quad (3.8)$$

- **The Aggregation Phase:** All messages arriving at node v_i from its neighborhood $\mathcal{N}(i)$ get combined. Because graphs have no fixed ordering or coordinate system, this step has to be permutation-invariant — the result can’t depend on what order the neighbors happen to be indexed in. Here, aggregation is guided by GAT attention coefficients α_{ij} , which weight chemically active neighbors more heavily than inert ones [52]:

$$\mathbf{M}_i^{(l)} = \sum_{j \in \mathcal{N}(i)} \alpha_{ij} \cdot \mathbf{W}_{\text{msg}}^{(l)} \cdot \mathbf{h}_j^{(l)}. \quad (3.9)$$

- **The Update Phase:** The aggregated message is combined with the node’s own current state to produce its representation for the next layer, passed through an ELU activation to give the model room to capture nonlinear structure [4, 11, 52]:

$$\mathbf{h}_i^{(l+1)} = \text{ELU}\left(\mathbf{M}_i^{(l)} + \mathbf{W}_{\text{self}}^{(l)} \cdot \mathbf{h}_i^{(l)}\right). \quad (3.10)$$

3.3.3 Multi-Label Mathematical Formulation:

Independent Sigmoids and Sparsity Mitigation

Since fragrance smells are complex and rarely just smell one thing, a mixture can read both smells at once. To handle this, the final layer skips a mutually-exclusive softmax and instead maps the 256-dimensional blend embedding through an MLP to L independent sigmoid outputs, one per target scent descriptor [37, 40, 45, 48]. Each output is a probability $y_{\text{pred},k}$ bounded between 0 and 1:

$$y_{\text{pred},k} = \sigma(x_k) = \frac{1}{1 + e^{-x_k}}. \quad (3.11)$$

A standard multi-label setup would optimize this with plain Binary Cross-Entropy, summing error across all L target classes independently [37]:

$$\mathcal{L}_{\text{BCE}} = -\frac{1}{K} \sum_{k=1}^L \left[y_k \log(y_{\text{pred},k}) + (1 - y_k) \log(1 - y_{\text{pred},k}) \right]. \quad (3.12)$$

However, olfactory data is severely imbalanced. Label density across the master dataset is typically below 3% [37], meaning a typical molecule or blend has only a few active descriptors out of the candidate vocabulary, with the vast majority sitting at zero. Across the whole matrix, that puts zeros at over 97% [3, 8]. Under plain BCE, those easy negatives dominate the gradient — a model can hit over 97% accuracy just by predicting zero everywhere, and in doing so never learns the rare positive labels that are actually informative.

To resolve this, the pipeline uses an Asymmetric Focal Loss, which splits the focusing behavior for positives and negatives into separate exponents, γ_p and γ_n , and adds a class-frequency weight α_k [17]:

$$\mathcal{L}_{\text{FL}} = -\frac{1}{K} \sum_{k=1}^L \left[y_k \cdot \alpha_k \cdot (1 - y_{\text{pred},k})^{\gamma_p} \cdot \log(y_{\text{pred},k}) + (1 - y_k) \cdot (1 - \alpha_k) \cdot y_{\text{pred},k}^{\gamma_n} \cdot \log(1 - y_{\text{pred},k}) \right]. \quad (3.13)$$

Setting $\gamma_p = 0$ leaves gradients from the rare positive labels untouched — they pass through at full strength. Setting $\gamma_n = 2$ does the opposite job for negatives: it exponentially suppresses the loss contribution from easy, correctly-predicted zeros, which is what strips away their gradient dominance. Together, these two settings push the model’s attention toward the rare positive notes instead of letting it coast on negatives, which is what makes reliable multi-label prediction possible on data this sparse [17].

Chapter 4

Methodology

This chapter explains the methodology adopted for designing, implementing and evaluating the structure-aware hierarchical graph neural network (SA-HGNN) proposed in this research work. First, the chapter discusses the overall research methodology, highlighting the variables (both independent and dependent) studied and controlled variables in this experiment. After that, it discusses the data collection and harmonization methods, which include the integration of Goodscents, Leffingwell and Sisson et al.’s [40] data sets into a combined dataset. Next, the chapter discusses the technical methodology, including the model and architecture, explaining the architecture configuration of Level 1 GAT encoder, Level 2 Set Transformer and classification layer that generate multi-label scent predictions. Then, the chapter tackles the out-of-distribution splitting and graph carving methodology used for generating a leakage-resilient test set. Finally, it discusses the metrics and statistical significance tests, and the ablation study design, used for analyzing the contribution of Level 2 Set Transformer.

4.1 Research Design

This study operates through a Quantitative Experimental Framework, employing supervised learning techniques on a multi-label molecular graph classification task. In this study, the independent variable is the Level 2 aggregation method, with three levels: the Set Transformer (the proposed method), mean pooling, and a reproduced pairwise blend model (the baselines). These are evaluated using four dependent variables: Hamming loss, macro F1, micro F1, and PR-AUC. The controlled

variables are the Level 1 GAT encoder, the data split, the loss function, the training hyperparameters, and the random seeds. By doing so, the effect of the aggregation mechanism is isolated in order to ensure that any difference in performance observed could not be attributed to anything but the interaction modeling technique.

4.2 Data Acquisition and Harmonization Pipeline

This section outlines the general procedures for collecting the raw single-molecule data from Leffingwell and GoodScents and blend-level data from Sisson et al.’s. It details how these data are merged, cleaned, and standardized to create one harmonized source.

4.2.1 Dataset Consolidation

Data acquisition will focus on three main datasets: Goodscents, Leffingwell, and Sisson et al.’s [2025]. The first two, accessible through Pyrfume, serve as the single-molecule sources for Level 1 pretraining and as the lookup reference for the molecular structures referenced in Sisson’s dataset. Sisson et al.’s dataset, accessible through their GitHub repository, serves as the reference for molecule blends used in Level 2 fine-tuning and evaluation.

Each of the two single-molecule datasets contains three files, namely: `behavior.csv`, which contains the stimulus and descriptor pairs data; `molecules.csv`, containing CID, MolecularWeight, IsomericSMILES, IUPACName, name; and the `stimuli.csv` file, which stores data such as Stimulus, TGSC ID, CID, Concentration %, Solvent.

The Leffingwell and Goodscents data files will be merged into one master dictionary, joined by the IsomericSMILES key present in both their `molecules.csv` files, where every molecule has both its canonical structure and its full set of odor descriptors attached. Every IsomericSMILES string in both sources is canonicalized with RDKit so that the same molecule written in different ways is matched correctly. Blend components in Sisson et al.’s dataset, which are identified by PubChem CIDs, are linked to the master dictionary by canonicalizing their SMILES in the same way. This was the key chosen as opposed to CID, as some of the CIDs contained in the files were negative/non-standard and would require further standardization, while the SMILES would already be canonicalized either way in the preparation stage for converting them into graphs. Furthermore, data such as the Concentration % and Solvent from the `stimuli.csv` files are to be dropped, since

concentration is scoped out from this study.

Descriptors from behavior.csv, which are indexed only by Stimulus, will first be joined to stimuli.csv on the shared Stimulus field to obtain each entry’s CID, which is then used to join to molecules.csv, which completes the within-source link between a compound’s structure and its descriptors before the cross-source SMILES merge described above takes place.

4.2.2 Vocabulary Standardization

To be more precise, there are approximately 669 unique raw string labels from Goodscents and 114 from Leffingwell, with “unique” here meaning distinct as literal text, since spelling variants of the same underlying descriptor (e.g., anisic/anise) are not yet consolidated at this stage. Upon closer inspection, there is a normalization problem arising wherein sources contain duplicate spellings or variants of scent descriptors, such as anisic/anise, wine/winey, cherry/cherry maraschino cherry, and caramel/caramellic. Manually reviewing each of the scents would be arduous at this scale; therefore, this study will adopt Sisson et al.’s own open-source canonicalization dictionary, which standardizes approximately 400 raw descriptors down to 104.

As the merged data may contain scent descriptors outside this dictionary, a fuzzy string-matching or edit-distance comparison using rapidfuzz’s Levenshtein-based similarity scoring will be done across the residual, uncovered portion of the combined raw label set to identify likely duplicate descriptors. Once identified, they will be manually reviewed and resolved, deciding whether each is mapped to an existing note or declared a genuinely new descriptor, with the latter still subject to being dropped if it does not appear across at least 30 molecules. Further normalization checks will be done, such as whitespace trimming and punctuation stripping.

This standardization applies to the single-molecule (Level 1 pretraining) vocabulary, which is permitted to exceed 104 notes if genuinely new descriptors survive the residual fuzzy-matching process described above. This is distinct from the blend-level (Level 2) classification target, which remains fixed at whatever the resulting count will be when the same survival criterion Sisson et al. applied to their carving run that narrowed down their 104-label vocabulary up to 74 is applied. This paper’s split may yield a similar but not necessarily identical count.

4.3 Technical Modeling and Architecture

This section presents the technical architecture of SA-HGNN in the order data flows through it: first Level 1, the Graph Attention Network that encodes each molecule independently from its atomic and bond-level features; then Level 2, the Set Transformer that allows the two resulting molecule embeddings to interact; and finally how the combined blend embedding is converted into per-descriptor predictions, along with the loss function used to train it.

4.3.1 Level 1 GAT Featurization

Once the dataset is read, the merged GoodScents and Leffingwell molecules are first clustered and split using UMAP and K-Means (Section 4.4.2), and only the molecules in the training clusters enter the Level 1 GAT featurization phase for standalone pretraining, so molecules held out for validation and testing are never seen during pretraining. Each molecule is represented as a graph with 46-dimensional node features and 10-dimensional edge features (Section 3.1.1, Tables 3.1–3.2). This GAT has three layers and uses LeakyReLU to compute raw attention scores between connected atoms, and normalizes these scores with the softmax formulation in Section 3.1.2.

Once the GAT processes a molecule’s graph, via a mean-pooling readout, its resulting 256-dimensional molecule vector passes through a Multi-Layer Perceptron that maps it to one output per standardized single-molecule descriptor, producing independent sigmoid outputs; the exact number of descriptors is reported after vocabulary standardization (Section 4.2.2).

Afterwards, a Binary Cross-Entropy loss function will be employed for backpropagation. This will be used as opposed to Focal Loss because the latter is more geared toward preventing models from becoming lazy and predicting all-negative outputs, a problem that worsens after Level 2’s graph-carving step. Focal Loss will instead be reserved for Level 2, described in Section 4.3.3.

After pretraining is completed, this temporary output head is discarded, leaving only a GAT with its encoder weights trained. In this step, the model is trained with the single-molecule dataset until it is ready for Level 2, wherein it enters a two-phase protocol for learning blend mechanics.

4.3.2 Level 2 Set Transformer Configuration

With Level 1 delivering a GAT trained on single molecules, the same GAT is applied to the component molecules of a blend pair, once per molecule, producing a 256-dimensional embedding for each:

$\mathbf{z}_A, \mathbf{z}_B \in \mathbb{R}^{256}$ Using the same GAT for both molecules ensures they are embedded under identical rules regardless of order. At this point, the two embeddings form an unordered set:

$$S_{\text{blend}} = \{\mathbf{z}_A, \mathbf{z}_B\} \subset \mathbb{R}^{256} \quad (4.1)$$

and it is the Set Transformer’s subsequent processing, rather than the shared GAT alone that guarantees the blend prediction itself is permutation-invariant: $\{\mathbf{z}_A, \mathbf{z}_B\} = \{\mathbf{z}_B, \mathbf{z}_A\}$ must yield the same result.

This set is passed through the Self-Attention Block (SAB), computing standard attention using Queries (Q), Keys (K), and Values (V) derived from the set itself:

$$\text{SAB}(X) = \text{MultiHead}(X, X, X) \quad (4.2)$$

As in the attention mechanism introduced in Section 3.1.2, each element attends to every other element in the set by comparing its Query against the Keys of all elements (including itself) to compute attention weights, which are then used to combine their Values into an updated representation. Because the same set X serves as the source of Q, K, and V simultaneously, this is self-attention rather than attention across two distinct sets. Concretely, for the two-element blend set, molecule A’s Query is compared against both A’s and B’s Keys, producing attention weights that combine A’s and B’s Values into an updated embedding for A with the same process occurring for B.

SAB is configured with 2 blocks and 4 attention heads, where the two embeddings are no longer treated as individual molecules but as components of a blend. Here, molecules A and B "learn" from each other via message passing, using 4 attention heads in parallel per block, across 2 sequential blocks, with each pass updating both embeddings based on the other’s influence. Formally, this yields the updated embeddings $\mathbf{z}'_A$ and $\mathbf{z}'_B$, where each is now a function of both original embeddings rather than of its own molecule alone. By the end of this process, $\mathbf{z}'_A$ and $\mathbf{z}'_B$ are no longer independent, having been shaped by each other’s influence, designed to allow the model to represent the receptor-level masking and competition discussed in Chapter 2.

Afterwards, $\{\mathbf{z}'_A, \mathbf{z}'_B\}$ undergoes Pooling by Multihead Attention (PMA), compressing the set into a single 256-dimensional blend embedding $\mathbf{z}_{\text{blend}} \in \mathbb{R}^{256}$. This compression uses attention instead of simple averaging so that it is possible for one molecule’s contribution to be weighted more heavily than the others depending on the learned relationship between them (e.g., a dominant, overpowering scent molecule masking a weaker one).

4.3.3 Output Head & Asymmetric Focal Loss:

The final blend embedding $\mathbf{z}_{\text{blend}} \in \mathbb{R}^{256}$ is passed through a new Multi-Layer Perceptron that contracts from $256 \rightarrow 128 \rightarrow L$, where L is the number of descriptors that survive the leakage-resistant OOD split described in Section 4.4, with up to 74, following Sisson et al.’s own carving outcome, with the exact value reported once this study’s own split is executed.. The output layer applies independent sigmoid activations across all L nodes, $x \in [0, 1]$, since scent descriptors are multi-label: each descriptor must be treated as an independent binary decision (1 = present, 0 = absent), unlike softmax, which forces all outputs to sum to 1.0, which is a formulation consistent with the multi-label framework established in Section 3.3.3. The resulting predictions are optimized using the Asymmetric Focal Loss defined in Section 3.3.3, with $\gamma_p = 0$ and $\gamma_n = 2$, following Ameta et al. [3].

As such, the final pipeline proceeds as follows for fine-tuning: single-molecule data is featurized and converted into a graph for Level 1; the same GAT trained on single molecules is then reused on those same molecules as components of a blend pair, aggregated together via SAB and PMA, and passed through the L -node sigmoid output head; predictions are scored using Asymmetric Focal Loss, with the resulting gradients updating weights across the entire network, including the pretrained Level 1 GAT end-to-end.

4.4 Leakage-Resistant OOD Splitting and Graph Carving

This section describes how the dataset is partitioned into training, validation, and test sets such that the resulting evaluation genuinely measures generalization to unseen chemistry, directly addressing the second research question. It explains the reason for why conventional scaffold split is unsuitable for this dataset, then describes the chemical-fingerprint clustering approach used instead, the graph-carving procedure that removes any blend pair spanning two partitions, the descriptor survival filter applied after carving, and finally a sanity check comparing this split against a standard random split.

4.4.1 Scaffold Split Decay

As discussed in more detail in Section 2.6, scaffold splitting does not work well for out-of-distribution (OOD) evaluation in sparse chemical spaces, where scaffold groups contain too few constituent

compounds, causing it to collapse to just random splitting instead. This step verifies this claim mathematically as opposed to just relying on the aforementioned literature. To do so, the molecule-to-scaffold ratio (R), is defined as:

$$R = \frac{N_{\text{molecules}}}{N_{\text{scaffolds}}} \quad (4.3)$$

where $N_{\text{molecules}}$ is the total number of unique canonicalized compounds, and $N_{\text{scaffolds}}$ stands for how many distinct Bemis–Murcko frameworks were generated using RDKit. As R gets closer to 1.0, it means that almost every scaffold only contains one molecule, meaning that there is virtually no shared molecular framework across the compounds. R will be computed based on the merged Goodscents + Leffingwell master dictionary during data prep and will be benchmarked against Fooladi et al.’s own collapse threshold of $R < 2.0$. Our R falling below theirs would further justify using the clustering-based alternative (Morgan-UMAP K-Means graph carving) detailed in the next section, rather than scaffold splitting.

4.4.2 Morgan-UMAP K-Means Split

To represent each molecule, ECFP4 fingerprinting will be performed via RDKit for all unique molecules in the dictionary, generating 2048 binary flags with a radius of 2, capturing atomic environments up to two bonds away.

These 2048-bit ECFP4 fingerprint vectors will then be mapped into a continuous 2D coordinate space using Uniform Manifold Approximation and Projection (UMAP), making it easier to cluster, with chemically and structurally similar molecules sitting closer together and vice versa. K-Means clustering will then be applied to the 2D map with $k = 7$. This value of $k = 7$ follows Guo et al. [14], who found seven UMAP clusters to outperform alternative clustering methods on chemical molecule data.

The OOD test-cluster selection rule works as follows: among the 7 candidate clusters, only those satisfying the requirement that every retained descriptor has at least one training instance and one test instance after graph-carving are considered eligible. Among these eligible candidates, the cluster selected is whichever one retains the largest number of usable blend pairs, as opposed to selecting based on raw cluster size. This mirrors Sisson et al.’s own justification for their chosen split ratio, which similarly prioritized maximizing usable data and label coverage over a fixed proportion.

A second eligible cluster, the one retaining the next-largest number of usable blend pairs, is withheld as the validation cluster, and the remaining clusters form the training set in the selected

test and validation clusters are withheld from training, establishing the single-molecule boundary required for the Graph-Carving procedure described in Section 4.4.3.

4.4.3 The Graph-Carving Protocol

The Sisson et al. [40] pairwise blend dataset is treated as a meta-graph in the construction of the graph-carving procedure, wherein each node represents a unique canonical single molecule, and each edge represents a documented blend between two such molecules.

Carving is then performed by examining each edge: if the two molecules an edge connects do not lie within the same partition (training, validation, or test), that edge is discarded. Edges are retained only if both molecules of a blend pair belong to the same partition to ensure there is no component overlap between any two partitions, preventing leakage. The validation set will be used for early stopping and will be the basis for hyperparameter tuning and the decision threshold for the sigmoid outputs. The test set will be used only once, in the final evaluation.

This carving mechanism mirrors the approach Sisson et al. [40] applied to their own random 50:50 partition, which discarded approximately 83,000 pairs to yield 44,000 training and 40,000 testing pairs. Nevertheless, since this thesis will group its chemical space using the Morgan-UMAP K-means clustering method mentioned in section above, instead of a randomly split 50:50 set, the exact pair count and the label count that will survive the pipeline will only be noted after it is executed.

4.4.4 Survival Vocabulary Filter

With graph carving discarding a significant amount of data from the dataset, the frequencies at which descriptors appear are altered. There is a need for a filter strictly targeting the blend-level target matrix, ensuring that every descriptor in the final multi-label target vector is sufficiently represented in its positive appearances on both sides of the carved split, for meaningful training and evaluation. This is determined by:

$$d \in \mathcal{D}_{\text{eval}} \iff c_{\text{train}}(d) \geq x \quad \text{and} \quad c_{\text{test}}(d) \geq x \quad (4.4)$$

wherein $c_{\text{train}}(d)$ is the number of instances descriptor d is positive in the carved training partition $\mathcal{E}_{\text{train}}$, and $c_{\text{test}}(d)$ is the number of instances descriptor d is positive in the test partition. x , where x , set to $x = 5$, meaning each retained descriptor must be positive in at least 5 blends on both

the training and test side. x is the minimum support threshold, where x dictates the minimum number of positive blends required on each side. A descriptor failing this condition on either side $c_{\text{train}}(d) < x$ or $c_{\text{test}}(d) < x$ is dropped from the classification target matrix. This emulates the filtering logic Sisson et al. [40] applied to their own carving run, where their final set of 74 emerged from applying this same survival criterion to their standardized 104-label vocabulary.

4.4.5 Interpolative Baseline Sanity Check

This step addresses the question of how much performance the model loses specifically because of the OOD splitting done, versus using a naive random splitting technique. To do so, the same single-molecule and blend datasets are split three ways (training, validation, and test) at random instead of by UMAP cluster, while still applying the same meta-graph carving and survival filtering processes. Everything else also stays the same, such as the model architecture, training hyperparameters, loss functions, and evaluation metrics.

The OOD performance decay Δ_{OOD} is calculated using:

$$\Delta_{\text{OOD}} = \text{Metric}_{\text{random}} - \text{Metric}_{\text{OOD}} \quad (4.5)$$

wherein a large gap ($\Delta_{\text{OOD}} \gg 0$) means that random splitting does cause performance inflation because structurally similar molecules can fall on both sides of a random split. A small gap ($\Delta_{\text{OOD}} \approx 0$) indicates that the proposed GNN framework is able to generalize meaningfully even when confronted with structurally unfamiliar chemical families it has never encountered during training. For Hamming loss, where lower is better, the difference is computed as OOD score minus random-split score, so that a positive Δ still indicates the OOD result is worse. This step is purely secondary, and the UMAP-clustered OOD split remains the study’s primary reported evaluation benchmark.

4.5 Quantitative Metrics and Metric Invariance

This section defines the metrics used to evaluate model performance and explains why each is necessary given the severe label sparsity of the blend-prediction task. It first examines the limitations of precision, recall, and F1 under this sparsity and motivates the use of macro-averaged F1 as the primary metric, then introduces Hamming loss as a complementary metric and defines it formally.

4.5.1 Multi-Label Metric Invariances

Let TP (true positive) be a descriptor correctly predicted as present, FP (false positive) a descriptor incorrectly predicted as present, FN (false negative) a present descriptor the model missed, and TN (true negative) a descriptor correctly predicted as absent. Because only about 2 to 4 of the L descriptors are positive in a typical blend, an estimated 95% of all label entries are negative. Because of this, a model can look like its accurate simply by saying a label is “absent” every time.. Precision, recall and F_1 are computed only from TP, FP and FN [34], so they ignore the many easy true negatives. A model that predicts no descriptors has no true positives and therefore receives an F_1 of zero, which exposes this failure. For this reason, macro-averaged F_1 , which gives every descriptor equal weight regardless of how common it is [34], serves as the primary metric. Micro-averaged F_1 pools all predictions and summarizes overall performance, and PR-AUC measures how well the model ranks true descriptors above false ones across all decision thresholds.

4.5.2 Hamming Loss Integration

Hamming loss is also reported, as it is a standard multi-label metric [34] that counts every incorrect label decision over the full $N \times L$ label matrix. Because it also counts the many easy true negatives, a model that predicts no descriptors would still obtain a low Hamming loss of about 0.03. Hamming loss is therefore interpreted alongside the all-negative and label-frequency reference baselines rather than on its own. This is calculated by the formula:

$$\text{HL} = \frac{1}{N \cdot L} \sum_{i=1}^N \sum_{j=1}^L \mathbf{1}(y_{ij} \neq \hat{y}_{ij}) \quad (4.6)$$

Where N is the total number of evaluated blend samples, L is the number of surviving target descriptors, up to 74 per Sisson et al.’s own carving outcome (Section 4.3.3), with the exact value reported once this study’s own carving is executed $\hat{y}_{ij} \in \{0, 1\}$ is the model’s binary prediction given a decision threshold, and $\mathbf{1}(\cdot)$ is the indicator function (returning 1 if true, 0 if false). Smaller values indicate higher accuracy ($\text{HL} \in [0, 1]$, with $\text{HL} = 0$ representing a perfect match).

4.6 Statistical Significance Testing

This section addresses the risk that the metrics gathered from evaluation might reflect random noise, such as seeding or splitting luck, rather than real performance differences, using inferential statistical

testing. Recent work cautions against relying on naive paired t -tests for comparing machine learning models, since the t -test is strongly affected by outliers and is not valid when resampled test sets are used [34]. Instead, this study will employ bootstrap resampling, as recommended by recent work, since it shows both whether one model is better than another and by how much [5]. The Set Transformer and the baseline models will be scored on the same resampled test indices, so that any difference reflects the models themselves rather than the test samples drawn. To account for seeding luck, each model will also be trained with several random seeds.

To do so, the fixed, held-out test predictions of N blend instances are taken, and N predictions are sampled with replacement 10,000 times, wherein both the Set Transformer and the baseline model are evaluated on the same resampled indices to ensure a paired comparison. For each bootstrap iteration b , the paired metric difference $\Delta(b)$ is computed by:

$$\Delta(b) = \text{Metric}_{\text{SetTransformer}}(b) - \text{Metric}_{\text{Baseline}}(b) \quad (4.7)$$

yielding a distribution of 10,000 paired differences that characterizes the variability of the Set Transformer’s advantage over the baseline. This procedure is repeated independently for each metric in the 4-metric suite from Section 4.5.2.

The significance threshold will be set at $\alpha = 0.05$ and 95% Empirical Confidence Intervals (CIs), which show the plausible range of the true underlying performance difference, will be extracted from the 2.5th and 97.5th percentiles of the 10,000 bootstrap iterations. Here, a CI that does not cross zero indicates a statistically significant positive gain at $\alpha = 0.05$. Since statistical significance alone does not show whether a difference is practically meaningful, the size of the effect is reported as the mean difference together with its 95% CI, expressed in the units of each metric.

This significance testing framework will be deployed for the purpose of comparing the Level 2 Set Transformer against naive Set Readouts (Mean Pooling) across the 4-metric suite from Section 4.5.2.

4.7 Ablation Study Design

This section describes the ablation and comparison studies used to isolate the contribution of each architectural component, which directly addresses the first and third research questions. It begins by replacing the Level 1 GAT encoder with a simpler Graph Convolutional Network to test the value of

attention at the atom level, then compares the proposed Set Transformer against mean pooling, and a reproduced pairwise blend model to test the value of attention at the molecule level, and finally examines whether the resulting attention weights align with documented receptor-level biology.

4.7.1 Intra-Molecule Encoder (Level 1 GCN Ablation)

This section details the ablation study choices that will be done on the model, where individual architectural components are swapped out to isolate their contribution to performance.

Specifically, only a partial ablation study will be conducted at Level 1, wherein the Level 1 GAT encoder is replaced with a Graph Convolutional Network (GCN). This is tested in combination with the winning Level 2 configuration discovered in the following subsection 4.7.2.

Level 1 undergoing only a partial ablation is a deliberate choice, as the encoder it uses is only secondary to the main question this thesis is asking, which is about how molecules in a blend interact with each other, being addressed by Level 2, not Level 1. Doing a full factorial ablation for Level 1 would only introduce computational cost and complexity and would not significantly advance the central research problem.

4.7.2 Inter-Molecule Aggregation (Level 2 Ablation)

This is the study’s main ablation test, to determine whether the use of explicit attention-based cross-molecule interaction modeling can meaningfully outperform naive linear pooling. Specifically, the architectural component being swapped out is the Level 2 Set Transformer (SAB + PMA), which will be compared against the mean pooling baseline, which calculates the unweighted arithmetic mean of the component vectors, assuming that the perception of a scent mixture is a linear combination of its single-molecule components, using the formula:

$$\mathbf{z}_{\text{blend}} = \frac{1}{M} \sum_{m=1}^M \mathbf{z}_m \quad (4.8)$$

In addition to the two pooling baselines above, a reproduced pairwise blend model is also included, following the GIN-based architecture Sisson et al. [40] describe as their point-of-comparison model: each molecule in the pair is encoded independently by a Graph Isomorphism Network (GIN), and the two resulting embeddings are concatenated, rather than pooled, before being passed through the same $256 \rightarrow 128 \rightarrow 74$ output head and Asymmetric Focal Loss used elsewhere in this study. Unlike

Sisson et al.’s original evaluation, which used a random 50:50 split, this reproduction is trained and evaluated on the same UMAP-clustered OOD split as the rest of this study’s models, so the comparison isolates the effect of the aggregation method rather than the effect of a different data split.

Everything else will remain constant across the ablation, using the same GAT from Level 1, the same dataset and splitting techniques, training setup, metric suite, and statistical protocol, to ensure that the differences that arise from the ablation are attributed directly to the swaps done. The results of this ablation, combined with the evaluation metrics and significance testing, will directly determine whether this thesis’s central hypothesis, that explicit attention-based interaction modeling meaningfully outperforms naive linear pooling for multi-label blend prediction, is supported.

4.7.3 Attention–Biology Alignment Analysis

To address the third research question, the Pooling by Multihead Attention (PMA) weights the trained Set Transformer assigned to each molecule in a blend are extracted from the test set. The primary test compares blends whose ground-truth labels exhibit suppression against blends without suppression, testing with a Mann–Whitney U test whether suppressing blends show a significantly larger imbalance between the two molecules’ PMA weights than non-suppressing blends. Because this test draws on every suppression case present in the test set rather than a small number of specific molecules, it does not depend on whether any single documented pair happens to survive the leakage-resistant split.

As a secondary, qualitative check, the test set is also searched for three molecule pairs with directly documented receptor-level antagonism from Section 2.4: β -ionone with sulfur-containing malodors [12], L-menthol with 2,3,5-trimethylpyrazine (TMP) [20], and Florhydral with isovaleric acid [21]. Where any of these specific pairs is present, the model’s behavior is reported as consistent with the documented mechanism if the antagonist molecule receives the higher PMA weight and the corresponding masked descriptor is absent from the blend’s predicted labels. Given the small number of candidate pairs, this check is illustrative rather than a basis for a statistical claim, and its absence from the test set does not undermine the primary Mann–Whitney result above.

This analysis overall is treated as exploratory rather than confirmatory: alignment between PMA weight asymmetry and documented antagonism indicates the model’s attention behavior is consistent with known receptor-level mechanisms, but does not constitute independent biological

validation of the underlying mechanism.

4.8 Sustainable Substitute-Finder Prototype

This section will describe the implementation of the sustainable substitute-finder prototype mentioned in Chapter 1 and Section 2.7, in support of UN SDG 12. Specifically, it will rely on the model’s Level 1 GAT encoder, because substitution is inherently a one-on-one single-molecule matching problem, making the Level 1 embedding space more appropriate for this task than the Level 2 space.

A queried SMILES molecule will go through the trained Level 1 GAT encoder, producing a 256-dimensional vector $\mathbf{z}_{\text{query}} \in \mathbb{R}^{256}$, which will then undergo a k -nearest-neighbor cosine similarity search against the pool of candidate substitute embeddings $\mathbf{z}_k$. Cosine similarity was chosen because it measures directional orientation in the latent space rather than vector magnitude, which is what determines qualitative semantic features in learned embedding spaces. This is consistent as well with methodology used by Tom et al. [2025] in their POMMix study. The distance metric formula is as follows:

$$\text{CosineSimilarity}(\mathbf{z}_{\text{query}}, \mathbf{z}_k) = \frac{\mathbf{z}_{\text{query}} \cdot \mathbf{z}_k}{\|\mathbf{z}_{\text{query}}\| \|\mathbf{z}_k\|} \quad (4.9)$$

The candidate pool for substitutes will consist of the single molecules from the Goodscents + Leffingwell single-molecule master dictionary. To ensure viability, this pool is programmatically filtered by cross-referencing the molecules’ SMILES strings and CIDs against published regulatory databases, specifically the IFRA Standards Library and EU REACH exclusion lists. Any compounds flagged as prohibited or restricted, marked as toxic/bioaccumulative, or sourced from endangered flora are permanently dropped from the search index prior to executing any queries.

In order to systematically assess the utility of this proof-of-concept without the need for wet-lab validation, the prototype’s recommendations will be evaluated using the following criteria:

- **Quantitative Evaluation (Precision@5):** The utility of the model will be measured by calculating the Precision at 5 (P@5) metric. This evaluates the percentage of the top five retrieved nearest neighbors that successfully share at least one standardized olfactory descriptor with the queried target molecule.
- **Regulatory and Safety Compliance (Constraint Satisfaction):** A recommendation is only practically useful to a fragrance chemist if it is legally and chemically safe. The prototype guarantees

this industrial utility by restricting its similarity search exclusively to the pre-screened candidate pool described above. By algorithmically masking flagged compounds from the vector space, it ensures that no IFRA-restricted, toxic, bioaccumulative, or endangered compounds can ever be recommended.

By satisfying these two criteria, the prototype demonstrates its applied utility as a sustainable, *in-silico* pre-screening tool. While full physical compatibility (such as Hansen solubility and ODT/OAV matching) and wet-lab confirmation are scoped out and deferred to future work, this structured evaluation ensures the model's outputs are both mathematically accurate and industrially viable.

Appendix A

Chemical Feature Vectors

The exact composition of the 46-dimensional node feature vector and the 10-dimensional edge feature vector used for the intra-molecular GAT encoder are detailed below.

Chemical Property	Dimensions	Mathematical / One-Hot Representation & Details
Atomic Identity	9	One-hot encoding of key organic and halogen elements in flavor/fragrance compounds: [C, N, O, S, F, Cl, Br, I, Other].
Formal Charge	1	Scalar representation of the localized net electrical charge on the atom.
Atomic Degree	6	One-hot representation of the number of connected covalent bonds to heavy atoms: [0, 1, 2, 3, 4, 5+].
Explicit Valence	7	One-hot encoding of the total valence electrons committed to bonding in the molecule: [0, 1, 2, 3, 4, 5, 6+].
Hybridization State	6	One-hot representation of the orbital configuration: [sp , sp^2 , sp^3 , sp^3d , sp^3d^2 , Other].
Aromaticity	1	Binary flag indicating if the atom is part of a resonance-stabilized aromatic ring structure.
Attached Hydrogens	5	One-hot encoding of the number of hydrogens bound to this heavy atom: [0, 1, 2, 3, 4+].
Ring Size Participation	6	Binary flags representing specific ring sizes to map cyclic functional groups: [3-ring, 4-ring, 5-ring, 6-ring, 7-ring, 8-ring].
Ring Membership	1	Binary flag indicating whether the atom is enclosed within any ring structure.
Chirality Tag	4	One-hot representation of tetrahedral stereocenter configuration: [Unspecified, Clockwise (CW), Counterclockwise (CCW), Other].

Table A.1: Specification of the 46-Dimensional Atom (Node) Feature Vector (X_i).

Chemical Property	Dimensions	Mathematical / One-Hot Representation & Details
Bond Type	4	One-hot representation of the bond’s electronic sharing structure: [Single, Double, Triple, Aromatic].
Conjugation	1	Binary flag indicating if the bond belongs to a continuous, overlapping pi-electron system.
Ring Membership	1	Binary flag indicating if the bond is closed within a ring structure.
Stereochemical Config	4	One-hot representation of spatial configuration around double bonds: [None, Stereo-Z, Stereo-E, Other].

Table A.2: Specification of the 10-Dimensional Bond (Edge) Feature Vector (E_{ij}).

REFERENCES

- [1] R. ACHEBOUCHE, A. TROMELIN, K. AUDOUZE, AND O. TABOUREAU, *Application of artificial intelligence to decode the relationships between smell, olfactory receptors and small molecules*, Scientific Reports, 12 (2022), p. 18817.
- [2] R. L. ALMEIDA, V. G. MALTAROLLO, AND F. G. F. COELHO, *Overcoming class imbalance in drug discovery problems: Graph neural networks and balancing approaches*, Journal of Molecular Graphics and Modelling, 126 (2024), p. 108627.
- [3] D. AMETA, S. KUMAR, R. MISHRA, L. BEHERA, A. CHAKRABORTY, AND T. SANDHAN, *Odor classification: Exploring feature performance and imbalanced data learning techniques*, PLOS ONE, 20 (2025), p. e0322514.
- [4] A. R. N. AOUICHAOU, F. FAN, S. S. MANSOURI, J. ABILDSKOV, AND G. SIN, *Combining group-contribution concept and graph neural networks toward interpretable molecular property models*, Journal of Chemical Information and Modeling, 63 (2023), pp. 725–744.
- [5] Y. BESTGEN, *Please, don't forget the difference and the confidence interval when seeking for the state-of-the-art status*, in Proceedings of the Thirteenth Language Resources and Evaluation Conference, N. Calzolari, F. Béchet, P. Blache, K. Choukri, C. Cieri, T. Declerck, S. Goggi, H. Isahara, B. Maegaard, J. Mariani, H. Mazo, J. Odijk, and S. Piperidis, eds., Marseille, France, June 2022, European Language Resources Association, pp. 5956–5962.
- [6] J. B. CASTRO, T. J. GOULD, R. PELLEGRINO, Z. LIANG, L. A. COLEMAN, F. PATEL, D. S. WALLACE, T. BHATNAGAR, J. D. MAINLAND, AND R. C. GERKIN, *Pyrfume: A window to the world's olfactory data*, bioRxiv, (2022).
- [7] C. CHEN, M. SHENG, H. YUAN, X. PAN, H. TIAN, AND X. LOU, *Perceptual interaction of aroma compounds and their effects on the enhancement of flavor quality in foods: a review*, Trends in Food Science & Technology, 165 (2025), p. 105315.
- [8] S. CHITHRANANDA, J. AMORES, AND K. K. YANG, *Mapping the combinatorial coding between olfactory receptors and perception with deep learning*, bioRxiv, (2026).
- [9] G. COUREAUD, T. THOMAS-DANGUIN, J.-C. SANDOZ, AND D. A. WILSON, *Biological constraints on configural odour mixture perception*, J Exp Biol, 225 (2022).
- [10] R. DATTA, J. NISTANE, A. SOSE, H. SAHU, AND R. RAMPRASAD, *Machine learning for green solvents: Assessment, selection and substitution*, Advanced Science, (2025).

- [11] H. FOOLADI, T. N. L. VU, M. MATHEA, AND J. KIRCHMAIR, *Evaluating machine learning models for molecular property prediction: Performance and robustness on out-of-distribution data*, Journal of Chemical Information and Modeling, 65 (2025), pp. 9871–9891.
- [12] Y. FUKUTANI, M. ABE, H. SAITO, R. EGUCHI, T. TAZAWA, C. A. DE MARCH, M. YOHDA, AND H. MATSUNAMI, *Antagonistic interactions between odorants alter human odor perception*, Current Biology, 33 (2023), pp. 2235–2245.e4. doi: 10.1016/j.cub.2023.04.072.
- [13] D. GASSER, S. SINGH, J. FREIHERR, AND A. GRASSKAMP, *Owsum: algorithmic odor prediction and insight into structure-odor relationships*, Journal of Cheminformatics, 15 (2023).
- [14] Q. GUO, S. HERNANDEZ-HERNANDEZ, AND P. J. BALLESTER, *Umap-based clustering split for rigorous evaluation of ai models for virtual screening on cancer cell lines**, Journal of Cheminformatics, 17 (2025), p. 94.
- [15] R. HAGAN, D. MOORE, F. MCGLONE, AND S. C. WALKER, *Olfactory scene analysis: Does analytical visual processing predict superior identification of component odours in a complex mixture?*, Perception, 54 (2025), pp. 408–430. PMID: 40165414.
- [16] S. HAN, H. FU, Y. WU, G. ZHAO, Z. SONG, F. HUANG, Z. ZHANG, S. LIU, AND W. ZHANG, *Himgnn: a novel hierarchical molecular graph representation learning framework for property prediction*, Briefings in Bioinformatics, 24 (2023), p. bbad305.
- [17] F. HATAMI AND M. MORADI, *Asymmetric focal loss improves graph neural network prediction of drug-drug interactions*, 2026.
- [18] Y. P. HENG, H. Y. LEE, J. W. CHONG, R. R. TAN, K. B. AVISO, AND N. G. CHEM-MANGATTUVALAPPIL, *Incorporating machine learning in computer-aided molecular design for fragrance molecules*, Processes, 10 (2022).
- [19] M. HLADIŠ, M. LALIS, S. FIORUCCI, AND J. TOPIN, *Matching receptor to odorant with protein language and graph neural networks*, in The Eleventh International Conference on Learning Representations, 2023.
- [20] B. HU, H. ZHENG, H. LIANG, H. LV, Q. XIAO, Y. JIANG, Q. LU, Y. NIU, H. WANG, AND Z. XIAO, *Elucidating odorant masking effects in food matrices through olfactory receptor-based profiling*, bioRxiv, (2026).
- [21] J. HUANG, J. LIN, R. YUENG, S. WU, L. SOLLA, AND T. ACREE, *Masking effects on iso-valeric acid recognition by sub-threshold odor mixture*, Chemical Senses, 48 (2023), p. bjad047.
- [22] R. JOERES, D. B. BLUMENTHAL, AND O. V. KALININA, *Data splitting to avoid information leakage with datasail*, Nature Communications, 16 (2025), p. 3337.
- [23] A. KENGKANNA AND M. OHUE, *Enhancing property and activity prediction and interpretation using multiple molecular graph representations with mmgx*, Communications Chemistry, 7 (2024), p. 74.
- [24] M. LALIS, M. HLADIŠ, S. A. KHALIL, L. BRIAND, S. FIORUCCI, AND J. TOPIN, *M2or: a database of olfactory receptor-odorant pairs for understanding the molecular mechanisms of olfaction*, Nucleic Acids Research, 52 (2024), pp. D1370–D1379.

- [25] B. K. LEE, E. J. MAYHEW, B. SANCHEZ-LENGELING, J. N. WEI, W. W. QIAN, K. A. LITTLE, M. ANDRES, B. B. NGUYEN, T. MOLOY, J. YASONIK, J. K. PARKER, R. C. GERKIN, J. D. MAINLAND, AND A. B. WILTSCHKO, *A principal odor map unifies diverse tasks in olfactory perception*, Science, 381 (2023), pp. 999–1006.
- [26] J. LEE, Y. LEE, J. KIM, A. R. KOSIOREK, S. CHOI, AND Y. W. TEH, *Set transformer: A framework for attention-based permutation-invariant neural networks*, 2019.
- [27] Q. LI, S. LI, J. LUO, AND H. YUAN, *Predictive models of odor contribution and thresholds for volatiles in identification of novel crop aroma compounds*, Metabolites, 15 (2025).
- [28] K. LÓPEZ-PÉREZ, J. F. AVELLANEDA-TAMAYO, L. CHEN, E. LÓPEZ-LÓPEZ, K. E. JUÁREZ-MERCADO, J. L. MEDINA-FRANCO, AND R. A. MIRANDA-QUINTANA, *Molecular similarity: Theory, applications, and perspectives*, Artificial Intelligence Chemistry, 2 (2024), p. 100077.
- [29] E. MAKRI, N. NAKIS, L. SISSON, G. MINSKY, L. TASSIULAS, V. SATARIFARD, AND N. A. CHRISTAKIS, *Benchmark for assessing olfactory perception of large language models*, 2026.
- [30] C. PATINHA CALDEIRA, R. FARCAL, C. MORETTI, L. MANCINI, H. RAUSCHER, K. RASMUSSEN, J. RIEGO SINTES, AND S. SALA, *Safe and sustainable by design chemicals and materials: Review of safety and sustainability dimensions, aspects, methods, indicators, and tools*, EUR 30991 EN, Joint Research Centre (JRC), European Commission, Luxembourg, 2022.
- [31] A. PAUL, *Dextrosinistral reading of smiles notation: Investigation into origin of non-sense code from string manipulations*, Digital Chemical Engineering, 15 (2025), p. 100222.
- [32] R. PELLEGRINO, E. J. MAYHEW, J. MARGOLIS, M. ANDRES, A. B. WILTSCHKO, R. C. GERKIN, AND J. D. MAINLAND, *Odors smell like their components: A linear framework for predicting olfactory mixture perception*, bioRxiv, (2026).
- [33] S. QIN, S. JIANG, J. LI, P. BALAPRAKASH, R. C. VAN LEHN, AND V. M. ZAVALA, *Capturing molecular interactions in graph neural networks: a case study in multi-component phase equilibrium*, Digital Discovery, 2 (2023), pp. 138–151.
- [34] O. RAINIO, J. TEUHO, AND R. KLÉN, *Evaluation metrics and statistical tests for machine learning*, Scientific Reports, 14 (2024), p. 6086.
- [35] B. C. L. RODRIGUES, V. V. SANTANA, S. MURINS, AND I. B. R. NOGUEIRA, *Molecule generation and optimization for efficient fragrance creation*, Industrial & Engineering Chemistry Research, 63 (2024), pp. 14480–14494.
- [36] S. ROMAGNY, G. COUREAUD, AND T. THOMAS-DANGUIN, *Developmental switch from elemental to configural processing of odor mixtures in humans*, Developmental Psychobiology, 68 (2026), p. e70188.
- [37] K. SAINI AND V. RAMANATHAN, *Predicting odor from molecular structure: a multi-label classification approach*, Scientific Reports, 12 (2022), p. 13863.
- [38] E. I. SANCHEZ MEDINA, S. LINKE, M. STOLL, AND K. SUNDMACHER, *Gibbs–helmholtz graph neural network: capturing the temperature dependency of activity coefficients at infinite dilution*, Digital Discovery, 2 (2023), pp. 781–798.

- [39] H. SHIMAKAWA, A. KUMADA, AND M. SATO, *Extrapolative prediction of small-data molecular property using quantum mechanics-assisted machine learning*, npj Computational Materials, 10 (2024), p. 11.
- [40] L. SISSON, A. A. BARSAINYAN, M. SHARMA, AND R. KUMAR, *Deep learning for odor prediction on aroma-chemical blends*, ACS Omega, 10 (2025), pp. 8980–8992.
- [41] L. J. S. SOBRAL, F. L. FEITOSA, J. A. M. KUNZ, P. K. DINIZ, J. F. AVELLANEDA-TAMAYO, J. L. MEDINA-FRANCO, A. G. FILHO, AND C. H. ANDRADE, *Odor-sight: An interpretable graph neural network web platform for prediction of odorant activity*, Journal of Chemical Information and Modeling, 66 (2026), pp. 9721–9728.
- [42] S. STESHIN, *Lo-hi: Practical ml drug discovery benchmark*, 2023.
- [43] J. SUH, Y. HONG, AND C. PARK, *A comparative study of machine learning models on molecular fingerprints for odor decoding*, Communications Chemistry, 8 (2025), p. 278.
- [44] M. TALHI, A. WOLF, AND A. MONOD, *Breaking symmetry bottlenecks in gnn readouts*, 2026.
- [45] A. N. TAREKEGN, M. GIACOBINI, AND K. MICHALAK, *A review of methods for imbalanced multi-label classification*, Pattern Recognition, 118 (2021), p. 107965.
- [46] G. TOM, C. T. SER, E. RAJAONSON, S. LO, H. S. PARK, B. LEE, AND B. SANCHEZ-LENDELING, *Does this smell the same? learning representations of olfactory mixtures using inductive biases*, Machine Learning: Science and Technology, 6 (2025).
- [47] J. WANG, Q. ZHANG, W. FAN, Q. SHI, J. MAO, J. XIE, G. CHAI, AND C. ZHANG, *Deciphering olfactory receptor binding mechanisms: a structural and dynamic perspective on olfactory receptors*, Frontiers in Molecular Biosciences, Volume 11 - 2024 (2025).
- [48] T. WEN, X. CAI, AND J. LI, *Graph neural networks vs. traditional qsar: A comprehensive comparison for multi-label molecular odor prediction*, Molecules, 30 (2025).
- [49] D. S. WIGH, J. M. GOODMAN, AND A. A. LAPKIN, *A review of molecular representation in the age of machine learning*, WIREs Computational Molecular Science, 12 (2022), p. e1603.
- [50] Y.-X. WU, X. WANG, A. ZHANG, X. HE, AND T.-S. CHUA, *Discovering invariant rationales for graph neural networks*, 2022.
- [51] Z. WU, J. WANG, H. DU, D. JIANG, Y. KANG, D. LI, P. PAN, Y. DENG, D. CAO, C.-Y. HSIEH, AND T. HOU, *Chemistry-intuitive explanation of graph neural networks for molecular property prediction with substructure masking*, Nature Communications, 14 (2023), p. 2585.
- [52] H. XIE, J. SUN, Y. SHAO, S. LI, S. HOU, Y. SUN, AND J. WANG, *Molecular odor prediction based on multi-feature graph attention networks*, 2025.
- [53] L. XU, D.-J. ZOU, AND S. FIRESTEIN, *Odor mixtures: A chord with silent notes*, Frontiers in Ecology and Evolution, Volume 11 - 2023 (2023).
- [54] M. ZAHEER, S. KOTTUR, S. RAVANBAKSH, B. POCZOS, R. SALAKHUTDINOV, AND A. SMOLA, *Deep sets*, 2018.

- [55] M. ZHANG, Y. HIKI, A. FUNAHASHI, AND T. J. KOBAYASHI, *A deep position-encoding model for predicting olfactory perception from molecular structures and electrostatics*, NPJ systems biology and applications, 10 (2024), p. 76.
- [56] W. ZHU, Y. ZHANG, D. ZHAO, J. XU, AND L. WANG, *Hignn: Hierarchical informative graph neural networks for molecular property prediction equipped with feature-wise attention*, 2022.